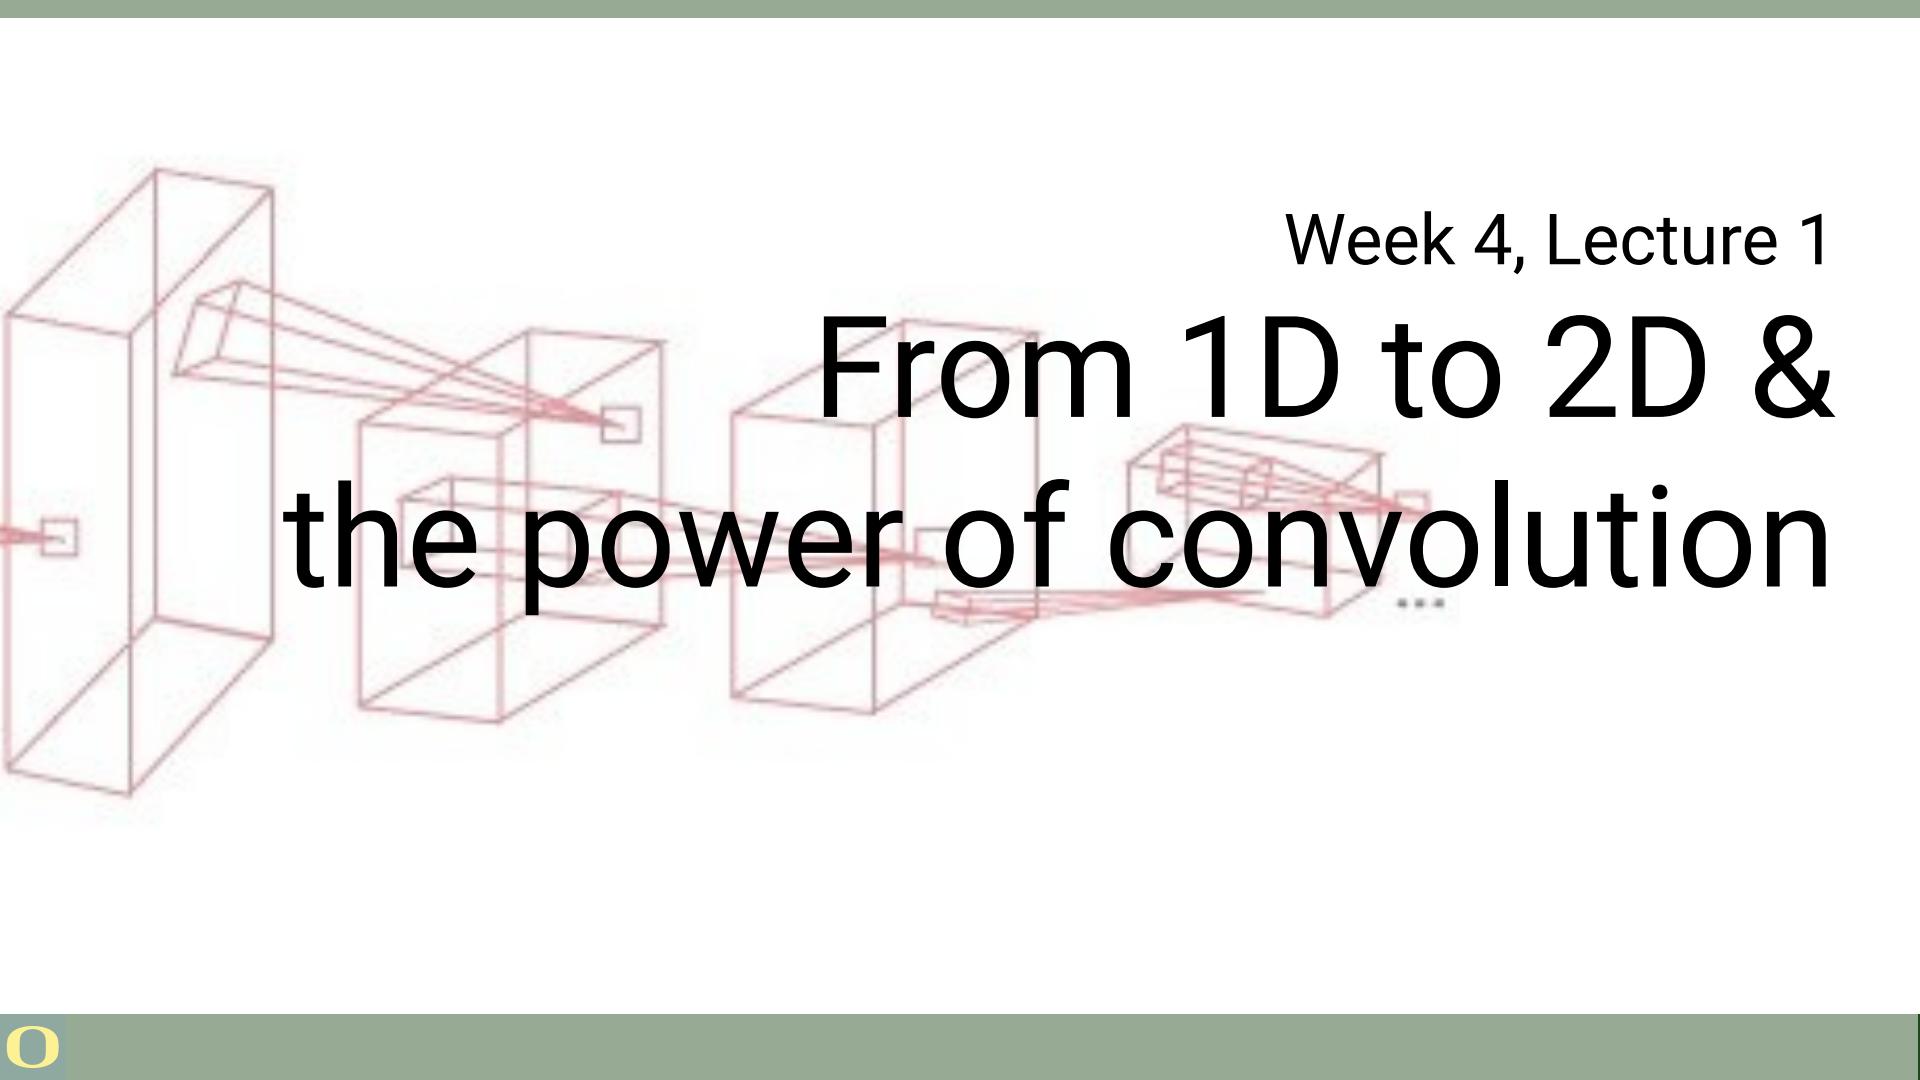


# Class core values

1. Be **respectful** to yourself and others
2. Be **confident** and believe in yourself
3. Always do your **best**
4. Be **cooperative**
5. Be **creative**
6. Have **fun**
7. Be **patient** with yourself while you learn
8. Don't be shy to **ask "stupid" questions**
9. Be **inclusive** and **accepting**



Week 4, Lecture 1

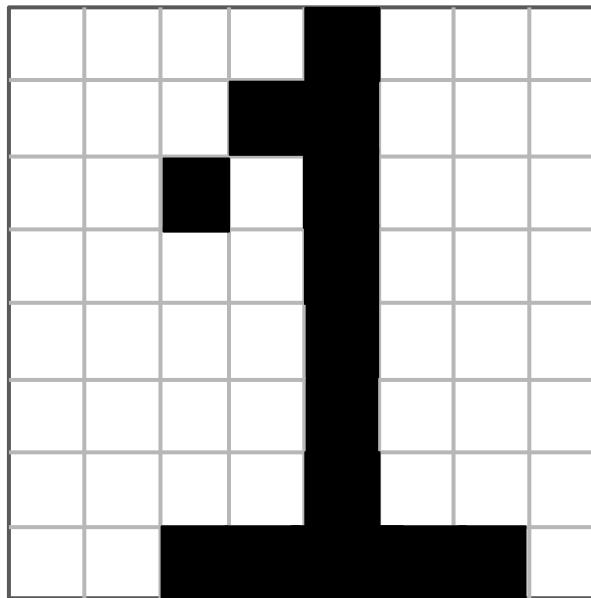
# From 1D to 2D & the power of convolution

# Learning Objectives

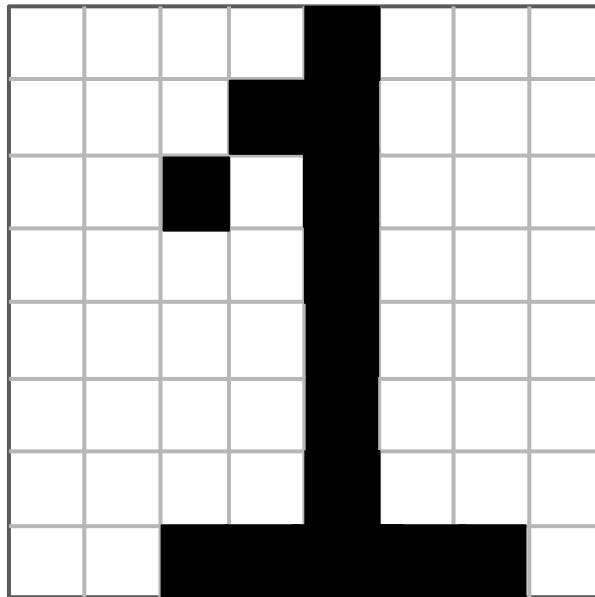
1. Describe the specific requirements of learning on image data
2. Explain the basic concept of convolution
3. Describe the specific terminology related to convolutional neural networks
4. Understand the applications of CNNs in image analysis

P1.

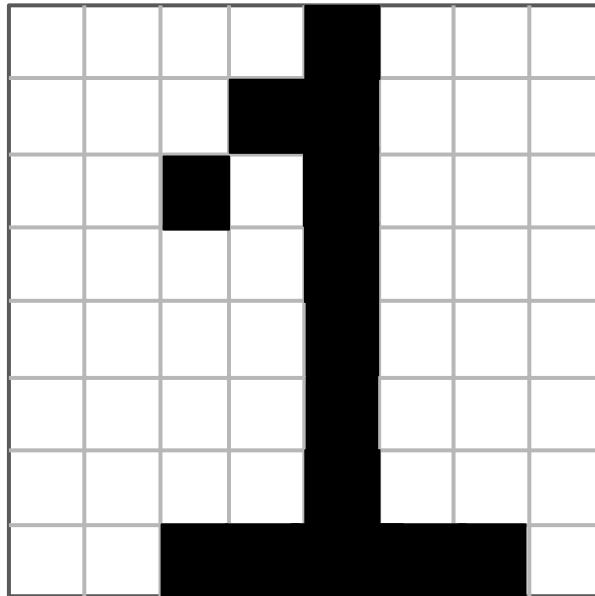
Can we use a dense neural net for image data?



# Images as input data



# Images as input data are matrices of numbers

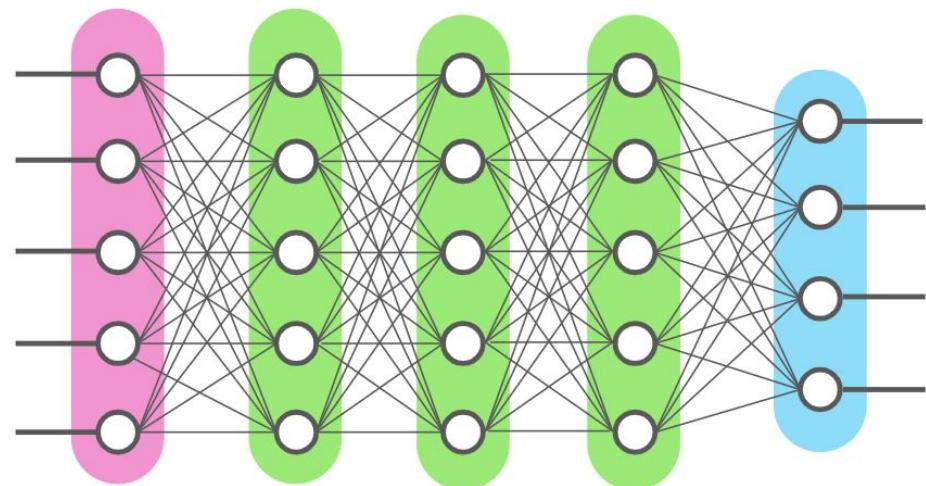


=

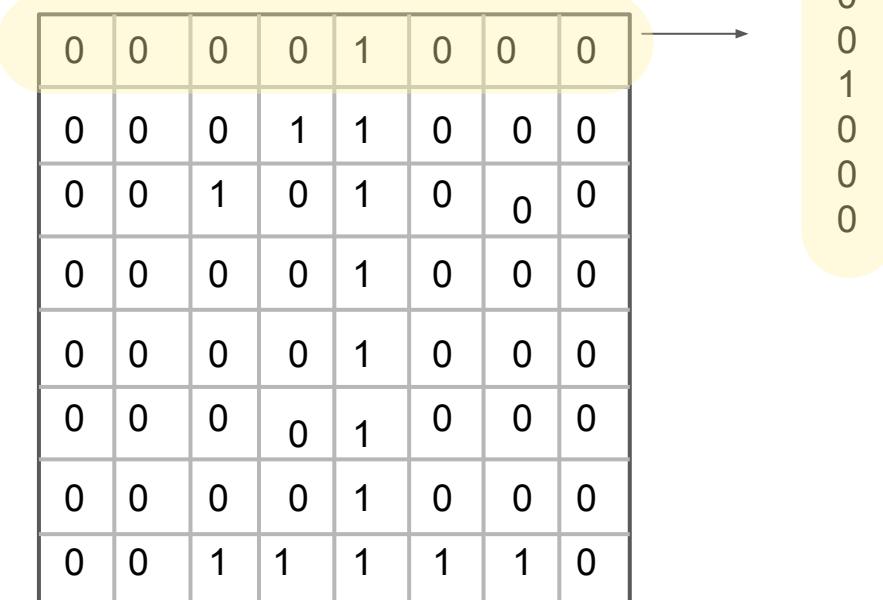
0	0	0	0	1	0	0	0
0	0	0	1	1	0	0	0
0	0	1	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	1	1	1	1	1	0

# Inputting image data to neural nets ?

0	0	0	0	1	0	0	0
0	0	0	1	1	0	0	0
0	0	1	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	1	1	1	1	1	0



# Inputting image data to neural nets ?



A 7x8 grid of binary values (0s and 1s) representing an image. The grid is highlighted with a yellow rounded rectangle. An arrow points from the bottom-right corner of the grid to a vertical vector on the right.

0	0	0	0	1	0	0	0
0	0	0	1	1	0	0	0
0	0	1	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	1	1	1	1	1	0

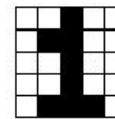
0  
0  
0  
1  
0  
0  
0

# Inputting image data to neural nets ?

0	0	0	0	1	0	0	0
0	0	0	1	1	0	0	0
0	0	1	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	1	1	1	1	1	0

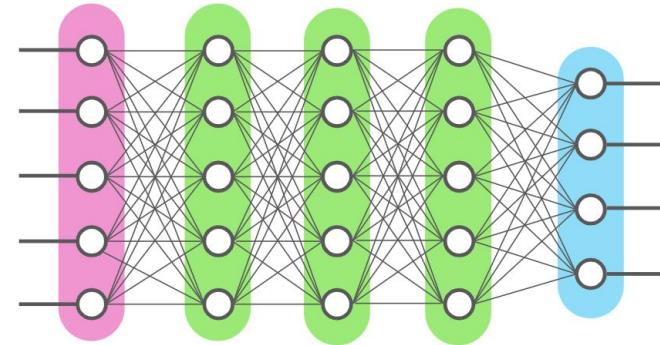
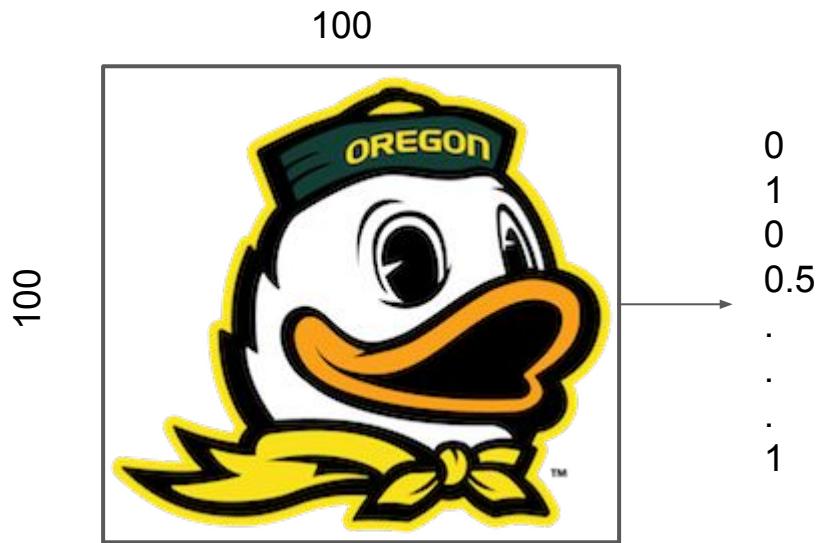


000100000011000

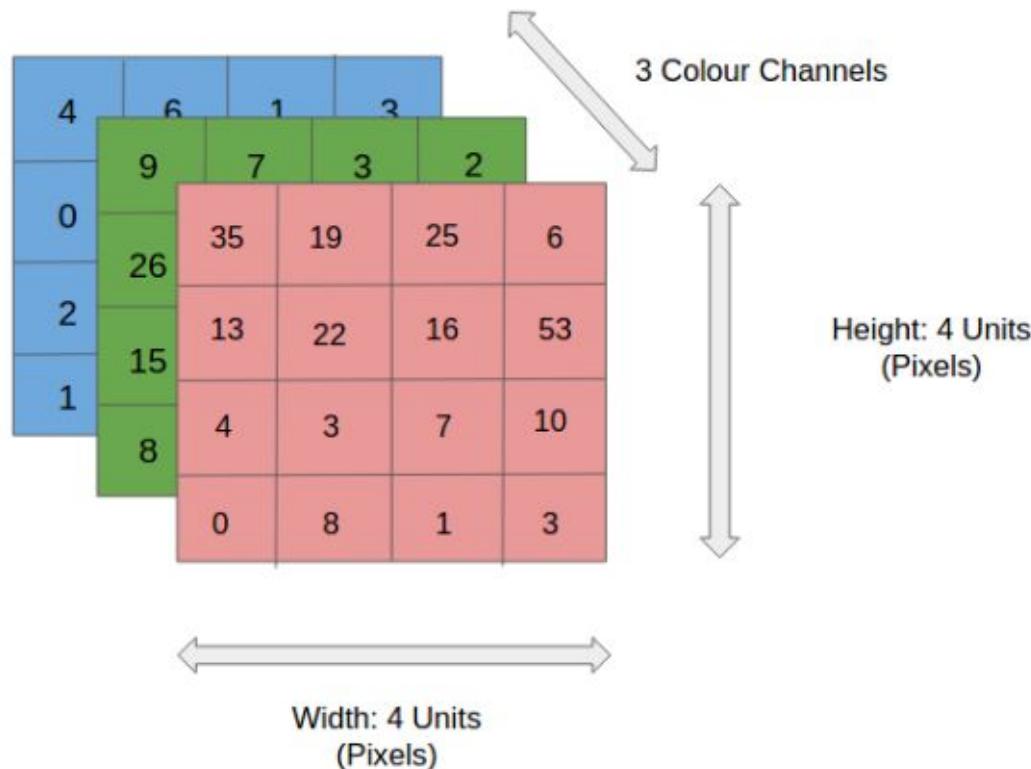


# In-class activity

How many connections?

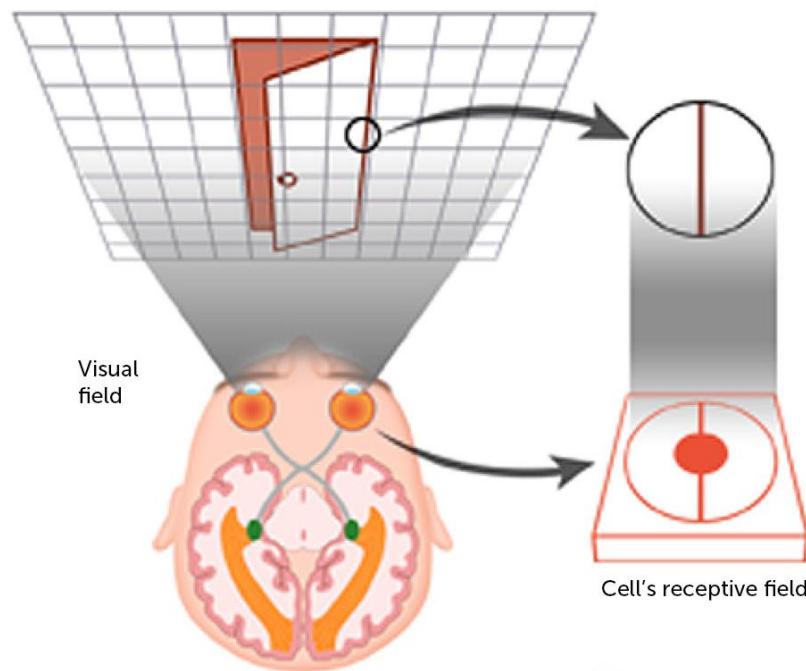


# Real images often have 3 channels with many pixels



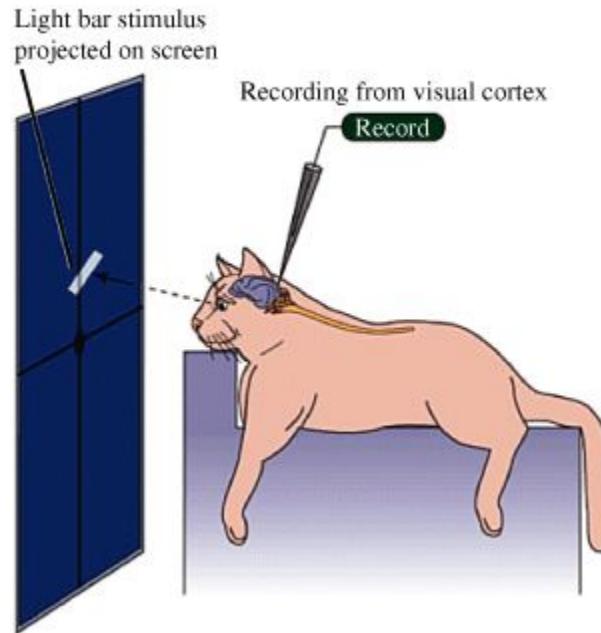
P2.

# Learning from our visual cortex

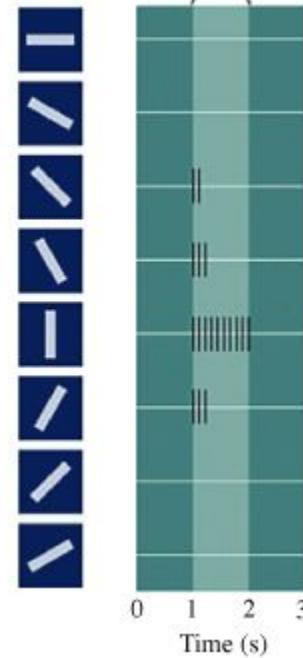


# Hubel and Wiesel's experiment on vision and the idea of receptive field

A Experimental setup



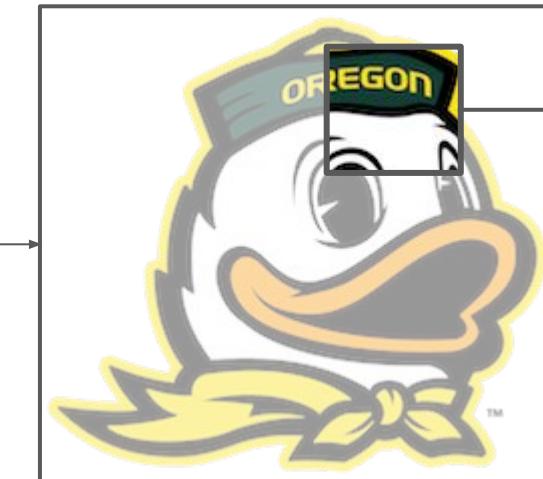
B Stimulus orientation Stimulus presented



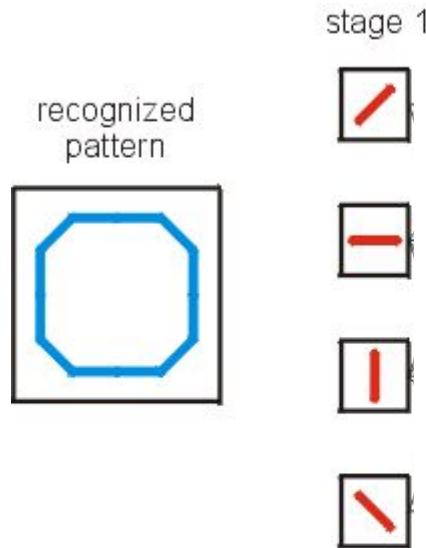
# Hubel and Wiesel's experiment on vision and the idea of receptive field



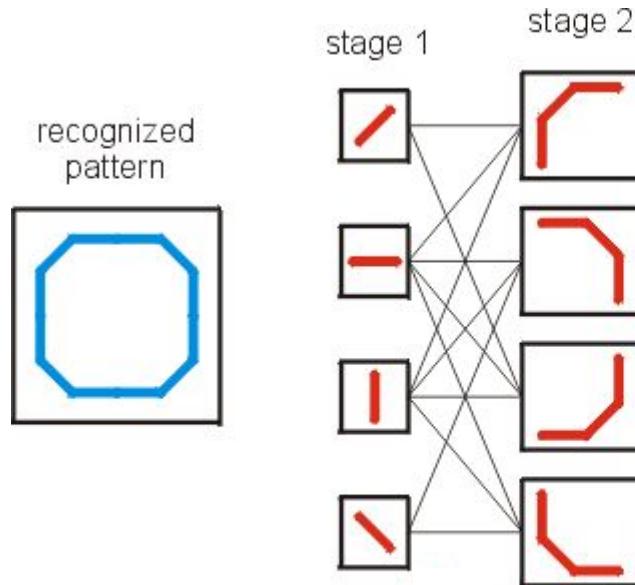
# Hubel and Wiesel's experiment on vision and the idea of receptive field



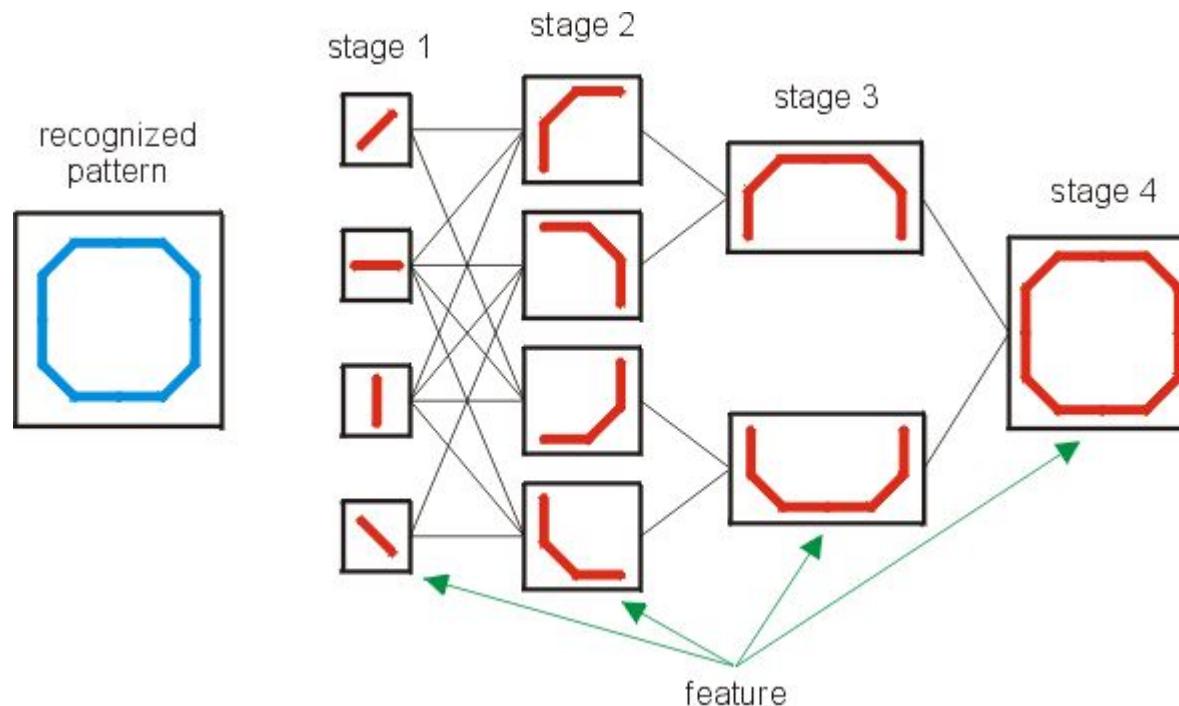
# Introduction of *neocognitron*: Features and feature extraction



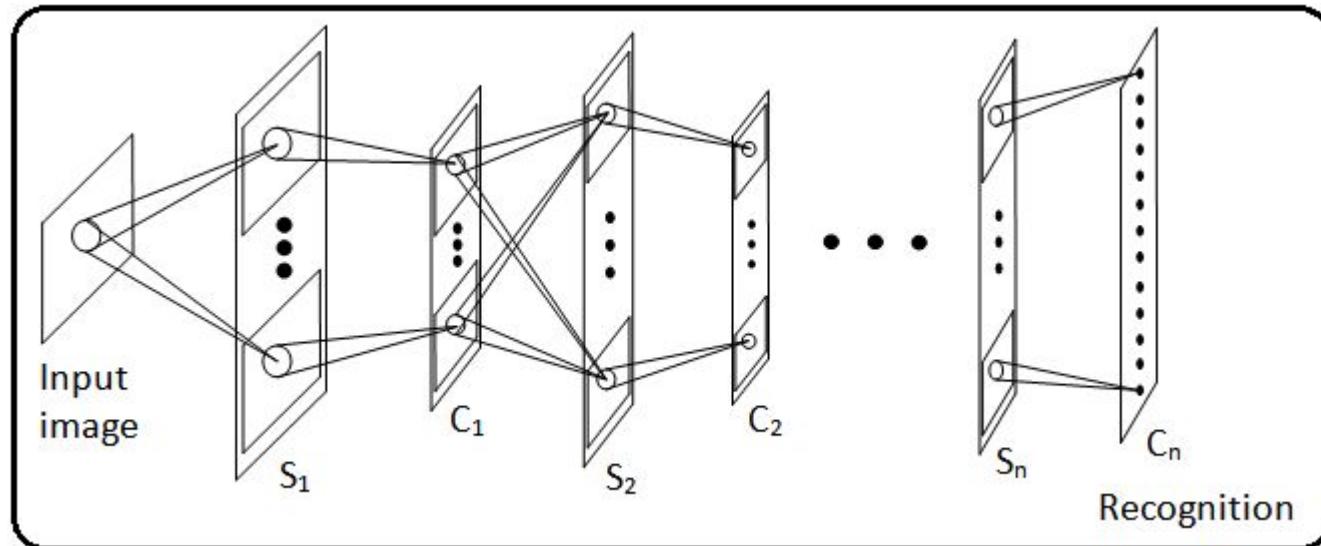
# Introduction of *neocognitron*: Features and feature extraction



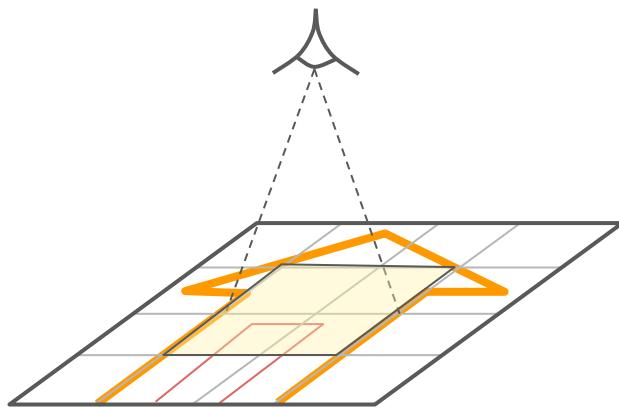
# Introduction of *neocognitron*: Features and feature extraction



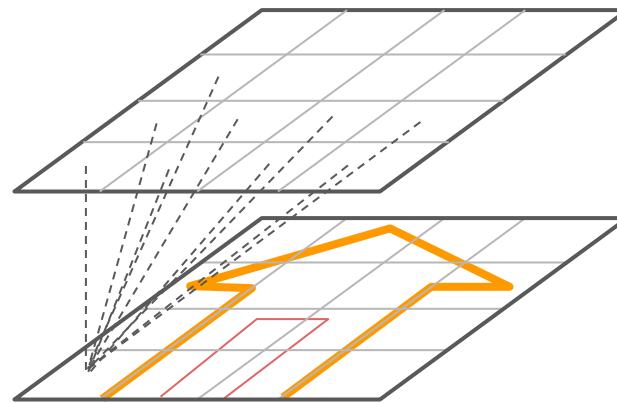
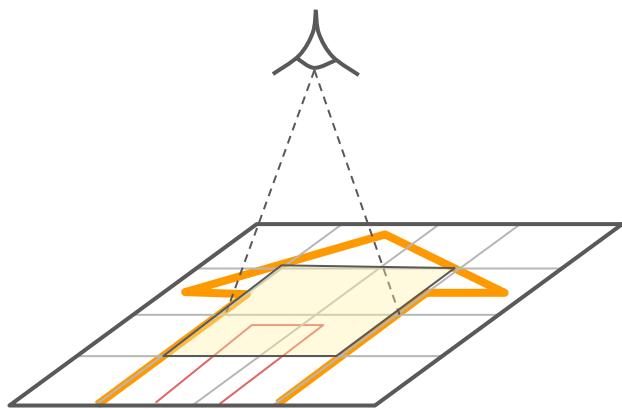
# Introduction of neocognitron



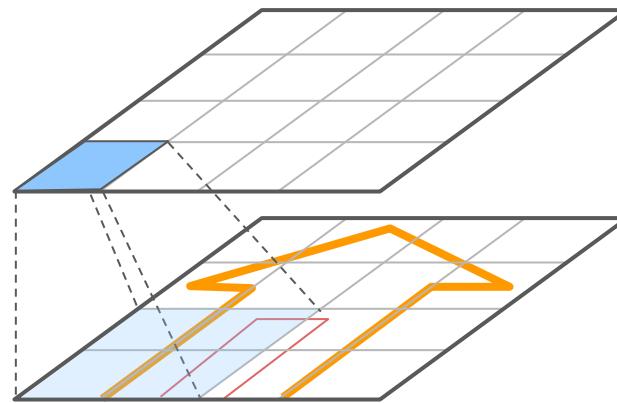
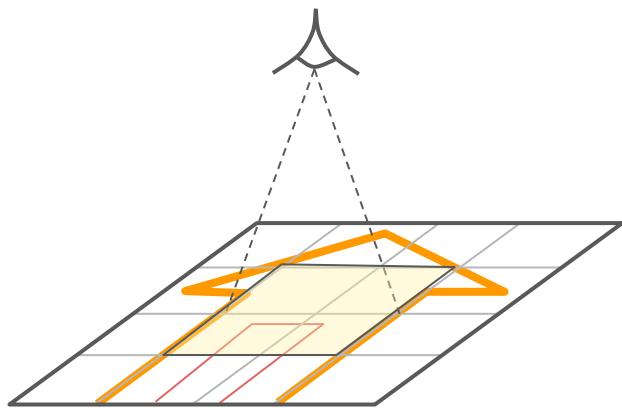
# Convolution, a breakthrough in image learning



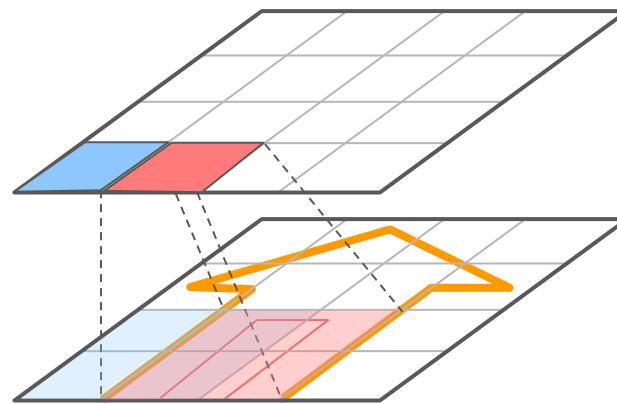
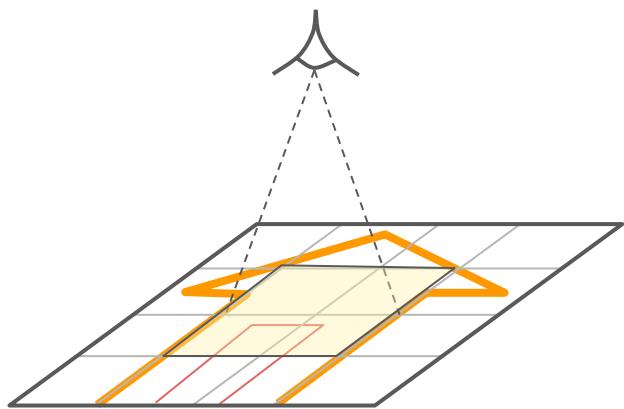
# Convolution, a breakthrough in image learning



# Convolution, a breakthrough in image learning



# Convolution, a breakthrough in image learning



Convolution is a mathematical calculation over the matrix that is similar to feature extraction

0	2	3	1	1	0
3	7	8	4	5	1
5	8	7	5	2	1
5	5	8	9	8	3
3	5	9	8	5	1
2	2	5	4	3	0

Convolution is a mathematical calculation over the matrix that is similar to feature extraction

0	2	3	1	1	0
3	7	8	4	5	1
5	8	7	5	2	1
5	5	8	9	8	3
3	5	9	8	5	1
2	2	5	4	3	0

0	1	0
0	1	0
0	1	0

Convolution is a mathematical calculation over the matrix that is similar to feature extraction

0x0	2x1	3x0	1	1	0
3x0	7x1	8x0	4	5	1
5x0	8x1	7x0	5	2	1
5	5	8	9	8	3
3	5	9	8	5	1
2	2	5	4	3	0

17

0	1	0
0	1	0
0	1	0

The amount of movement of the convolutional filter over the image is called *stride*

0	2x0	3x1	1x0	1	0
3	7x0	8x1	4x0	5	1
5	8x0	7x1	5x0	2	1
5	5	8	9	8	3
3	5	9	8	5	1
2	2	5	4	3	0

17	18
----	----

Stride = 1

# The amount of movement of the convolutional filter over the image is called *stride*

0	2x0	3x1	1x0	1	0
3	7x0	8x1	4x0	5	1
5	8x0	7x1	5x0	2	1
5	5	8	9	8	3
3	5	9	8	5	1
2	2	5	4	3	0

Stride = 1  
Stride = 2

0	2	3x0	1x1	1x0	0
3	7	8x0	4x1	5x0	1
5	8	7x0	5x1	2x0	1
5	5	8	9	8	3
3	5	9	8	5	1
2	2	5	4	3	0

# Convolution, a breakthrough in image learning

1 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	0	0
0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	1	0
0 <small><math>\times 1</math></small>	0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1	1
0	0	1	1	0
0	1	1	0	0

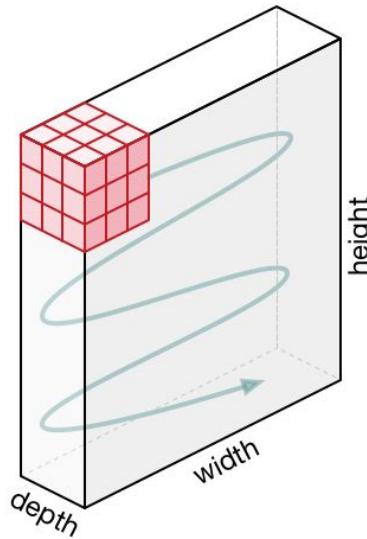
Image

$$m - n + 1$$

4		

Convolved  
Feature

# Convolution in reality is a 3D operation



0	0	0	0	0	0	0	...
0	156	155	156	158	158	158	...
0	153	154	157	159	159	159	...
0	149	151	155	158	159	159	...
0	146	146	149	153	158	158	...
0	145	143	143	148	158	158	...
...	...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	169	...
0	164	165	168	170	170	170	...
0	160	162	166	169	170	170	...
0	156	156	159	163	168	168	...
0	155	153	153	158	168	168	...
0	154	152	152	157	167	167	...
...	...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	165	...
0	160	161	164	166	166	166	...
0	156	158	162	165	166	166	...
0	155	155	158	162	167	167	...
0	154	152	152	157	167	167	...
...	...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



308

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-498

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



164

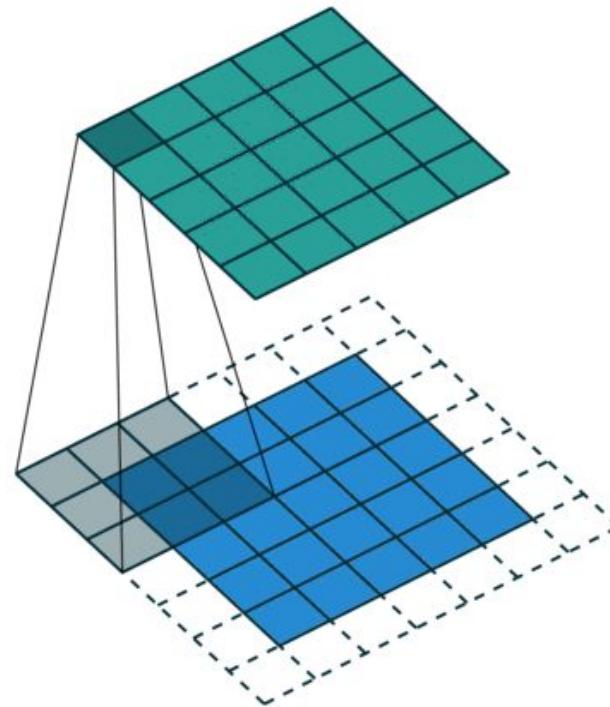
+ 1 = -25

Bias = 1

-25				...
				...
				...
				...
...	...	...	...	...

# In-class activity

Perform a simple convolution



*Padding gives more room to the filter and retains the size after convolution*

0	0	0	0	0	0	0	0	0
0	0	2	3	1	1	0	0	0
0	3	7	8	4	5	1	0	0
0	5	8	7	5	2	1	0	0
0	5	5	8	9	8	3	0	0
0	3	5	9	8	5	1	0	0
0	2	2	5	4	3	0	0	0
0	0	0	0	0	0	0	0	0

“Same” padding

vs

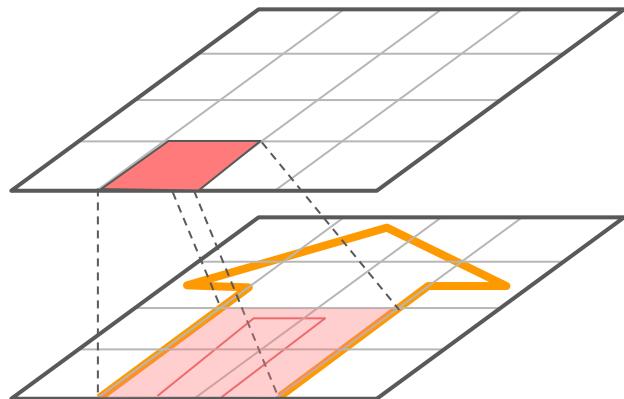
no padding  
(aka “valid” padding)

*Padding gives more room to the filter and retains the size after convolution*

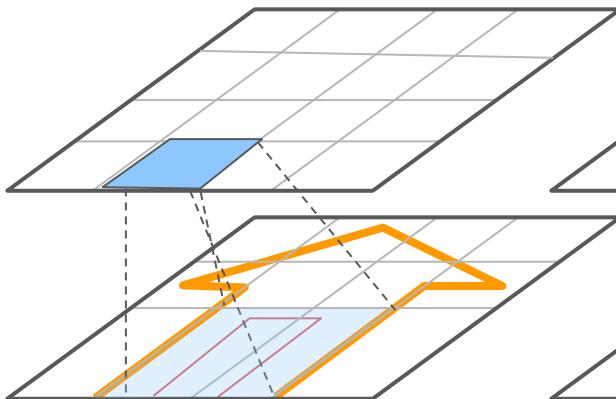
0	0	0	0	0	0	0	0	0
0	0	2	3	1	1	0	0	0
0	3	7	8	4	5	1	0	0
0	5	8	7	5	2	1	0	0
0	5	5	8	9	8	3	0	0
0	3	5	9	8	5	1	0	0
0	2	2	5	4	3	0	0	0
0	0	0	0	0	0	0	0	0

# Each convolution is a filter/kernel that generates one feature map

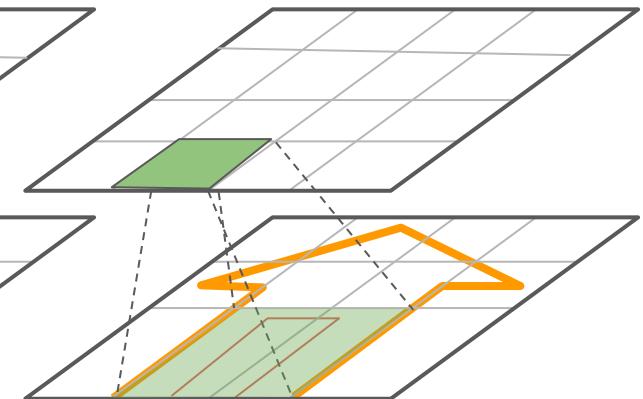
Feature map 1



Feature map 2

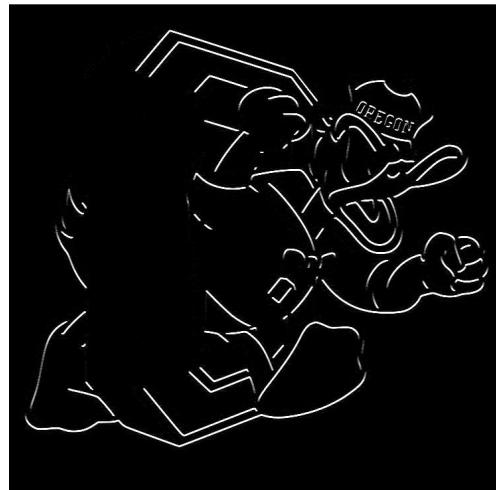


Feature map 3

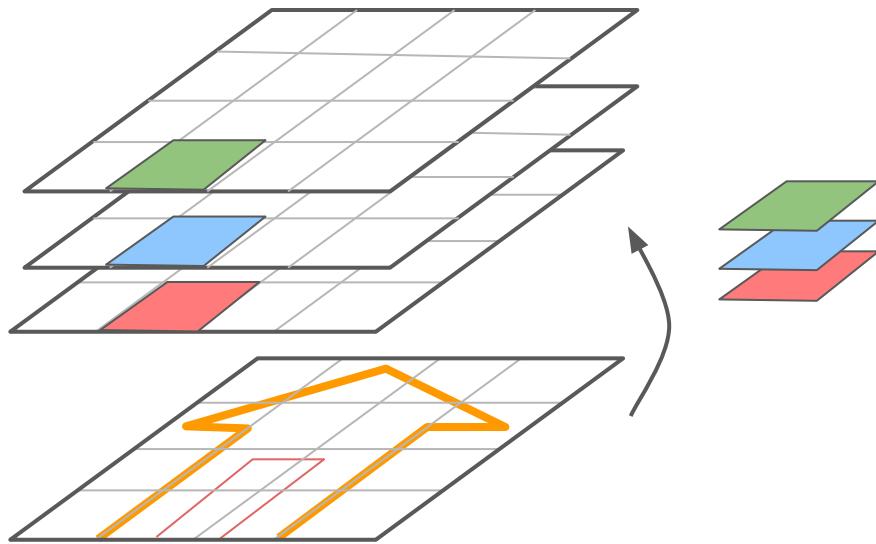


# In-class activity

## Visualizing the effect of kernels

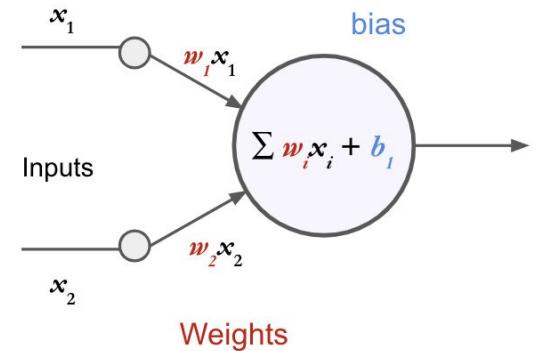
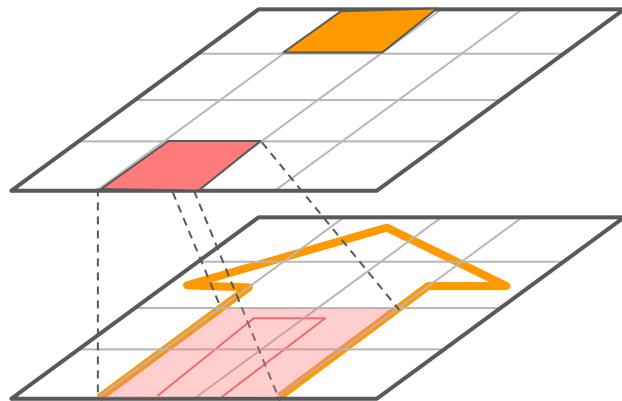


Each convolution is a filter/kernel that generates one feature map



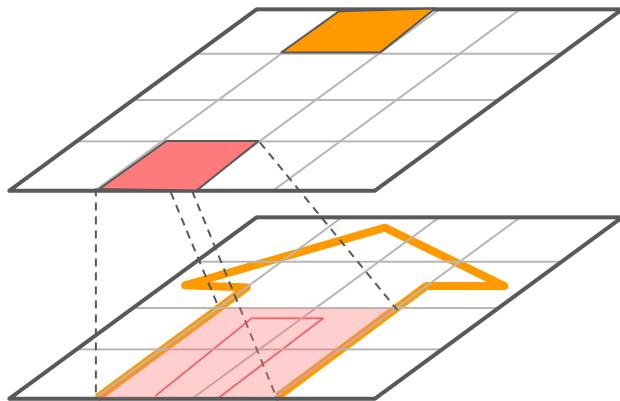
# Convolution filters are weights in a convolutional layer

Feature map 1



All neurons in one feature map share same parameters, aka are generated with same kernel

Feature map 1



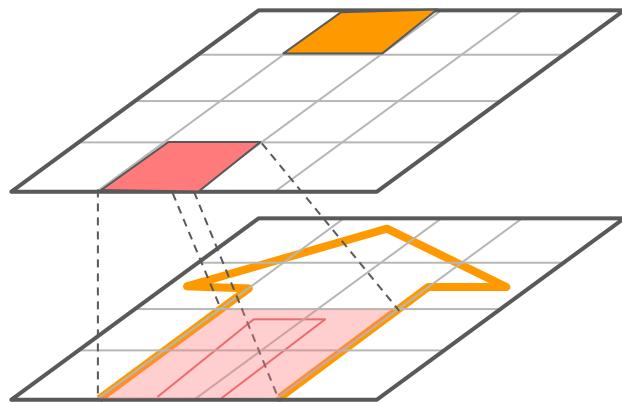
Same weights:

- Weights = convolution filter

Same bias

# Weight sharing enhances efficiency and learning

Feature map 1

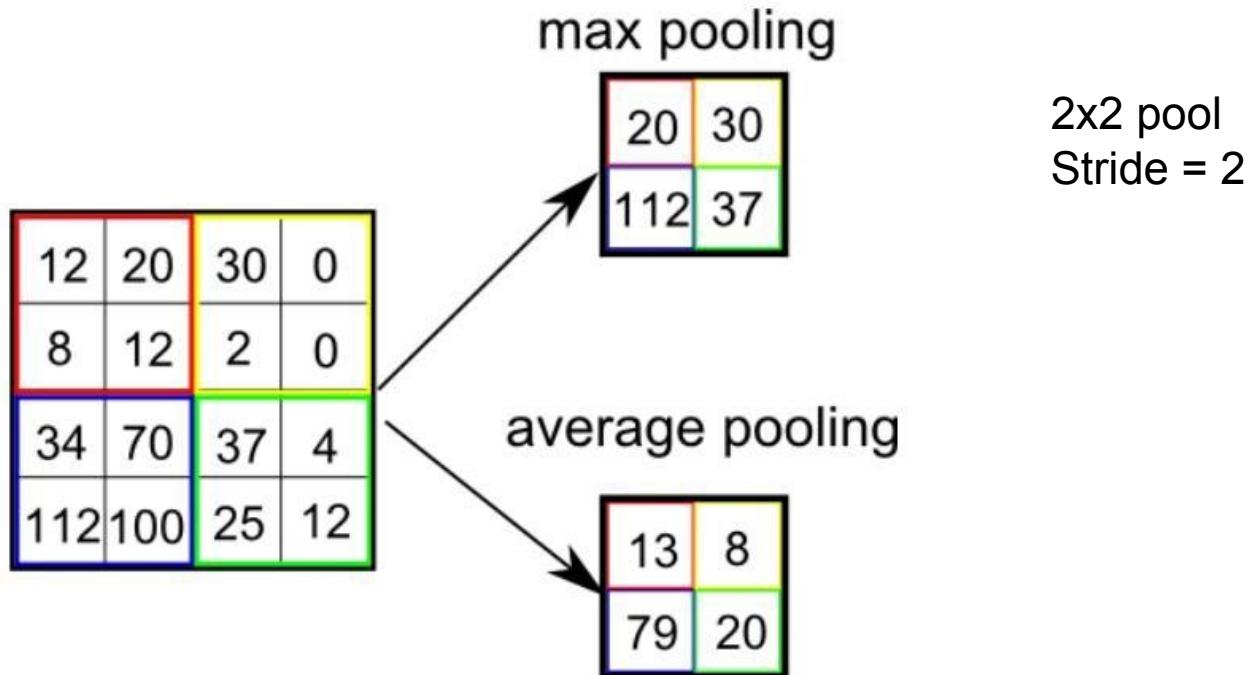


Same weights:

- Weights = convolution filter

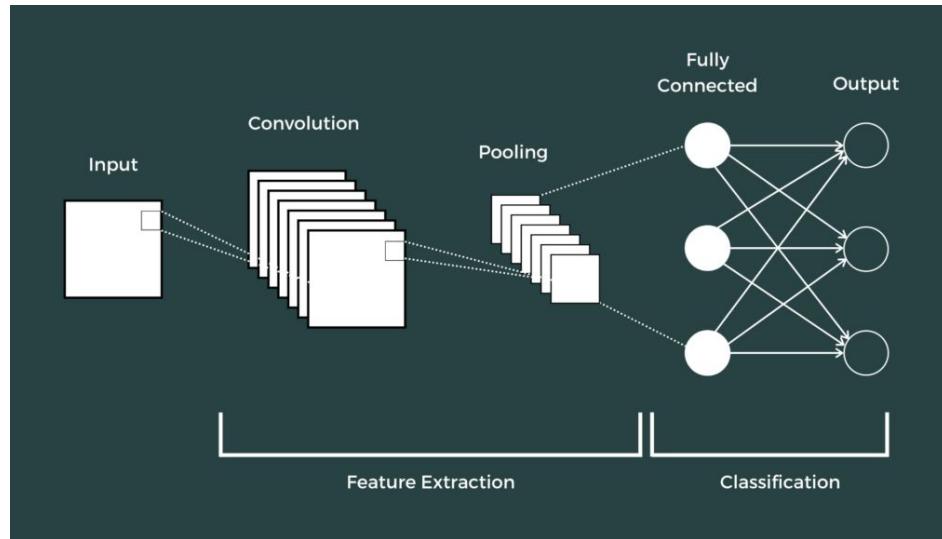
Same bias

*Pooling* layers allow for reduction of the convolved features



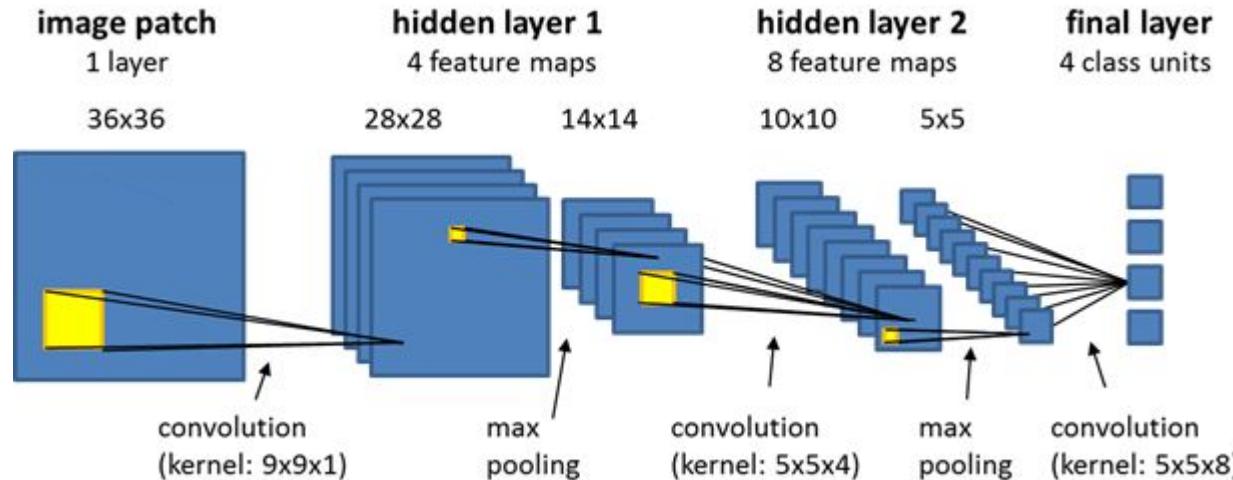
# In-class activity

## Shape of layers and architecture of model



# In-class activity

## Shape of layers and architecture of model



# Translation/Rotation and the concept of invariance



CNNs are not truly invariant to rotation or translation (equivariant to translation)

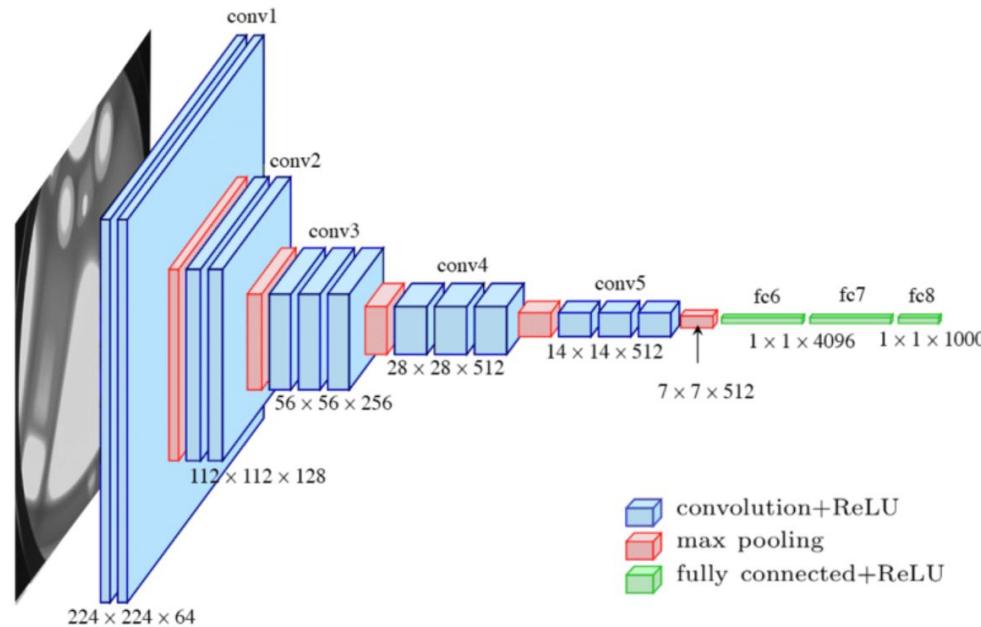


# Data augmentation is necessary



P4.

# Some common CNN architectures



# The power of a curated dataset and competitions



14,197,122 images, 21841 synsets indexed

[Home](#) [Download](#) [Challenges](#) [About](#)

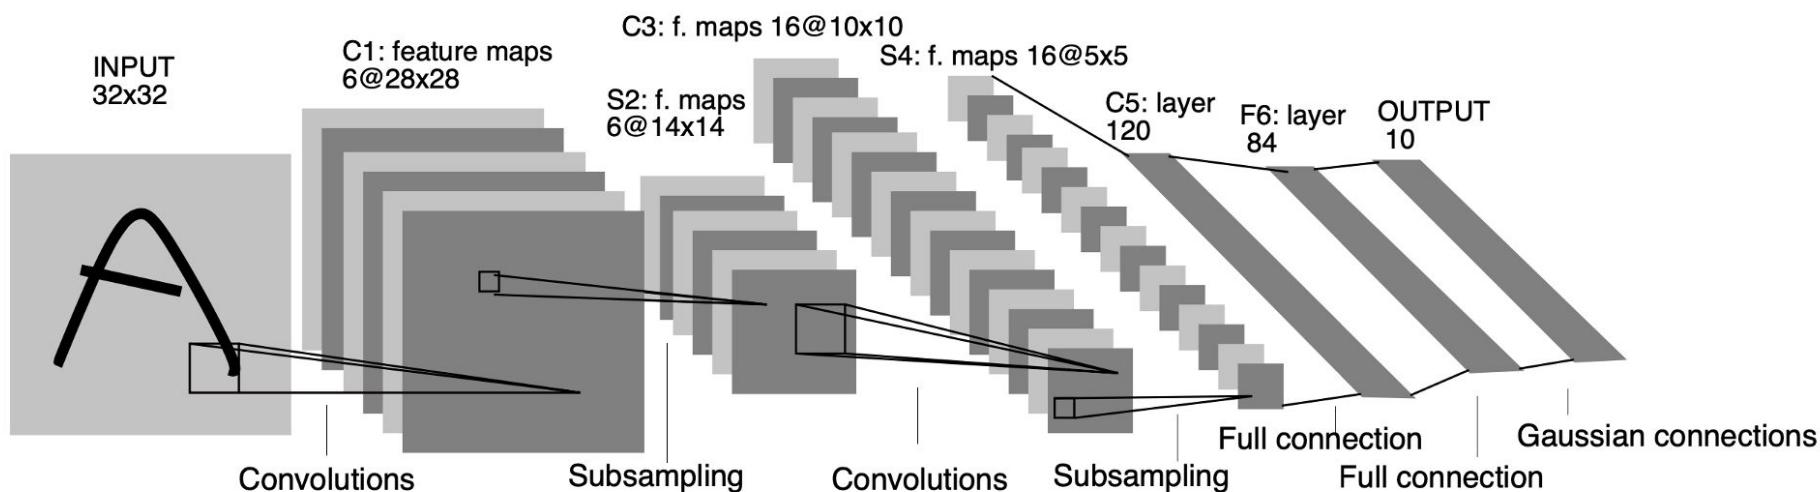
Not logged in. [Login](#) | [Signup](#)

**ImageNet** is an image database organized according to the **WordNet** hierarchy (currently only the nouns), in which each node of the hierarchy is depicted by hundreds and thousands of images. The project has been **instrumental** in advancing computer vision and deep learning research. The data is available for free to researchers for non-commercial use.

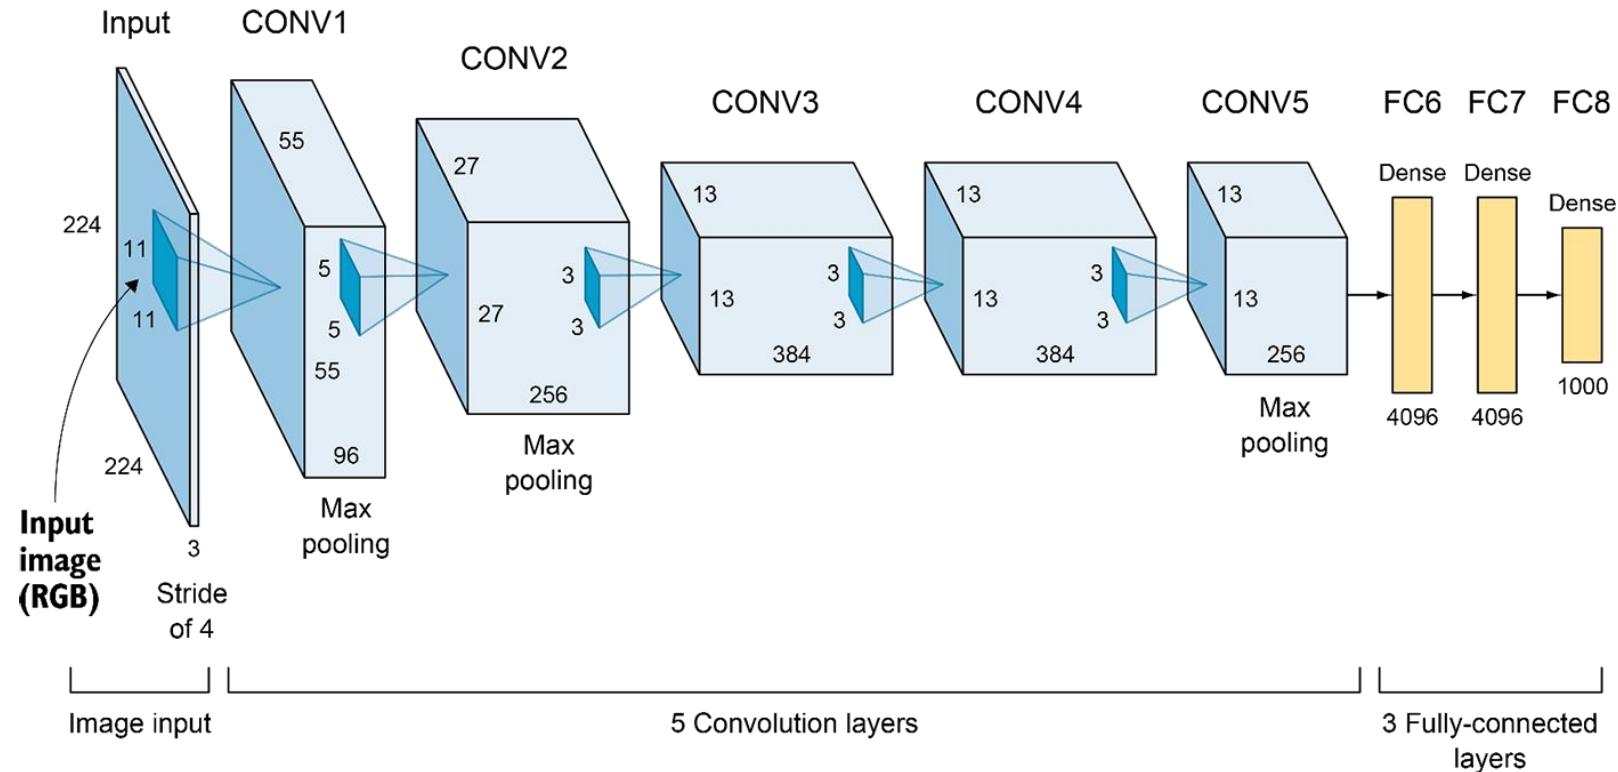
# The power of a curated dataset and competitions



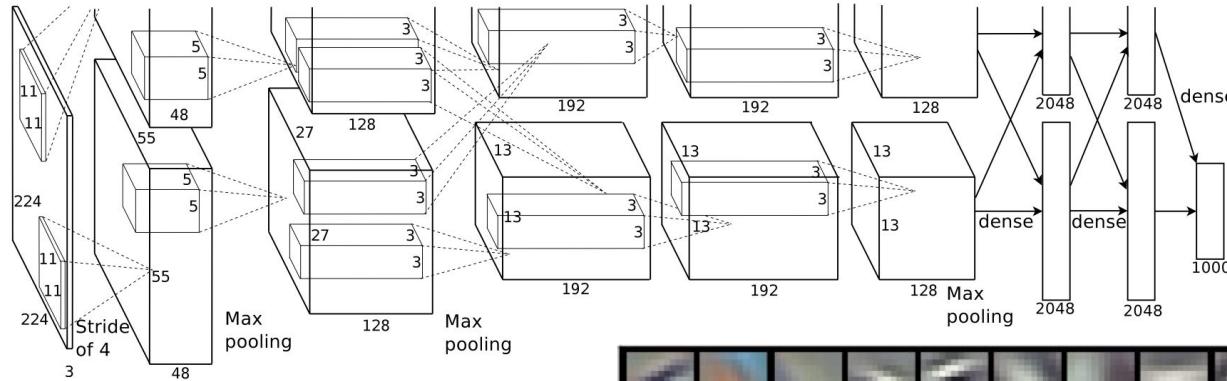
# LeNet-5



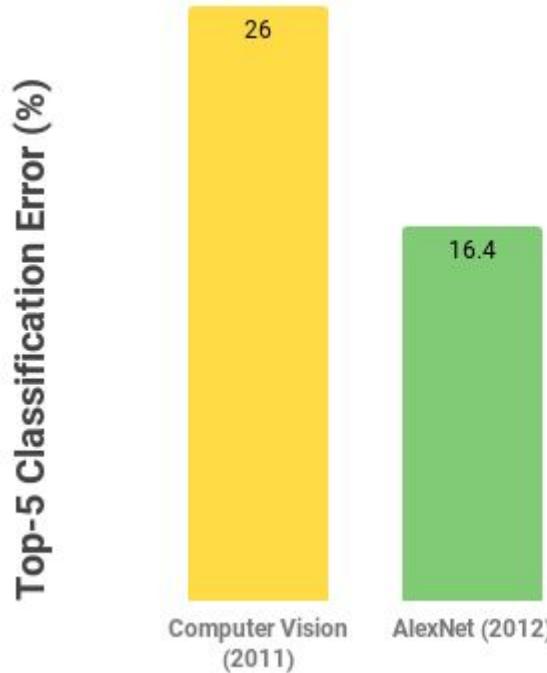
# AlexNet



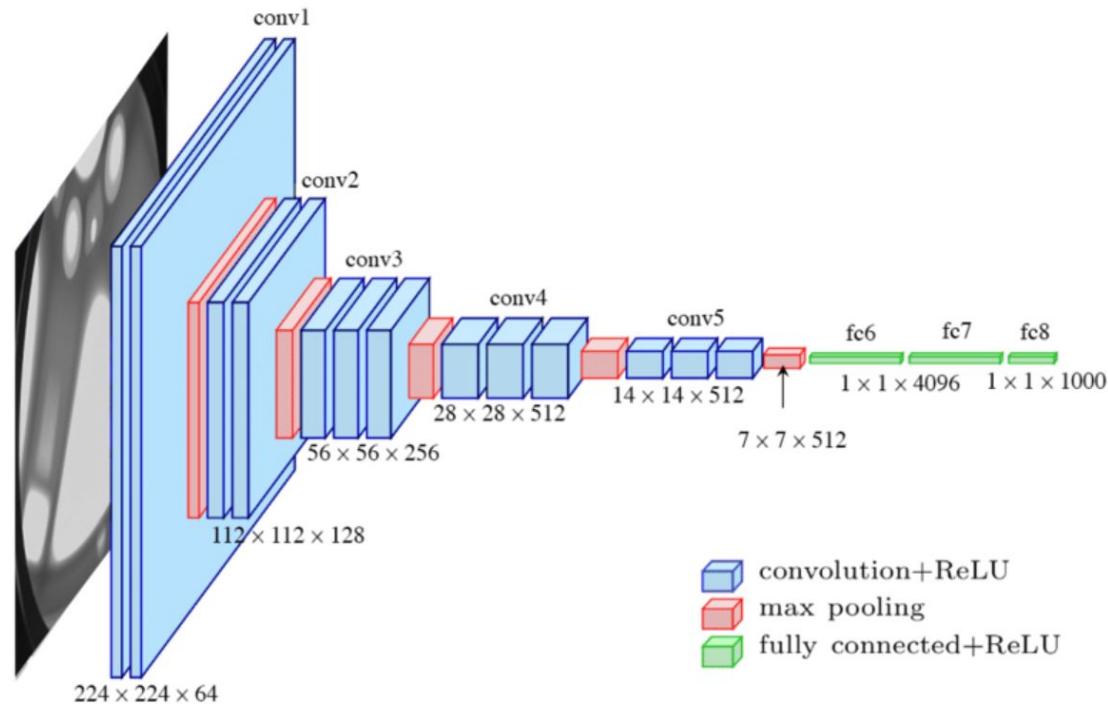
# AlexNet



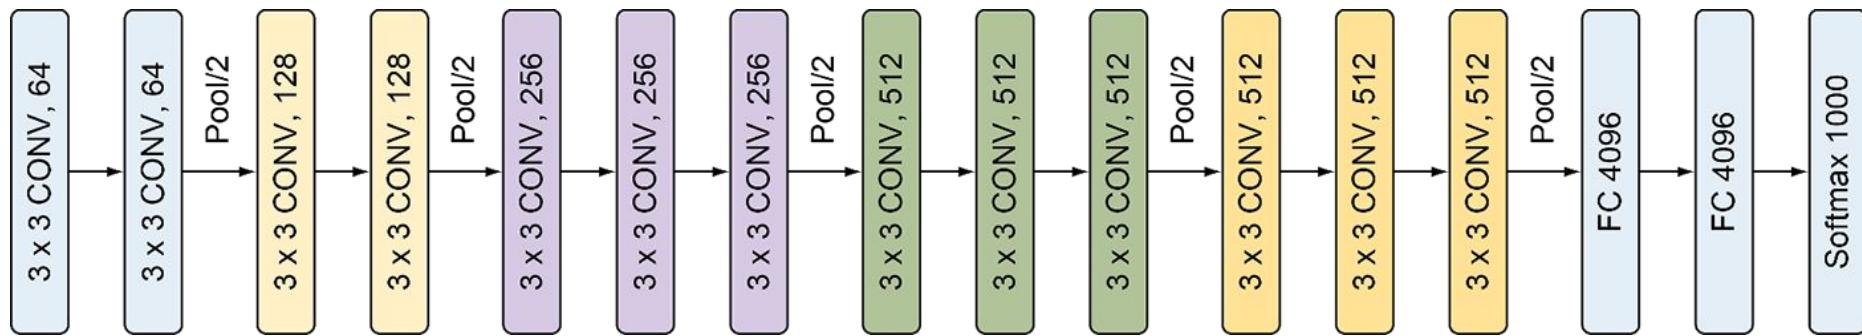
# The power of a curated dataset and competitions



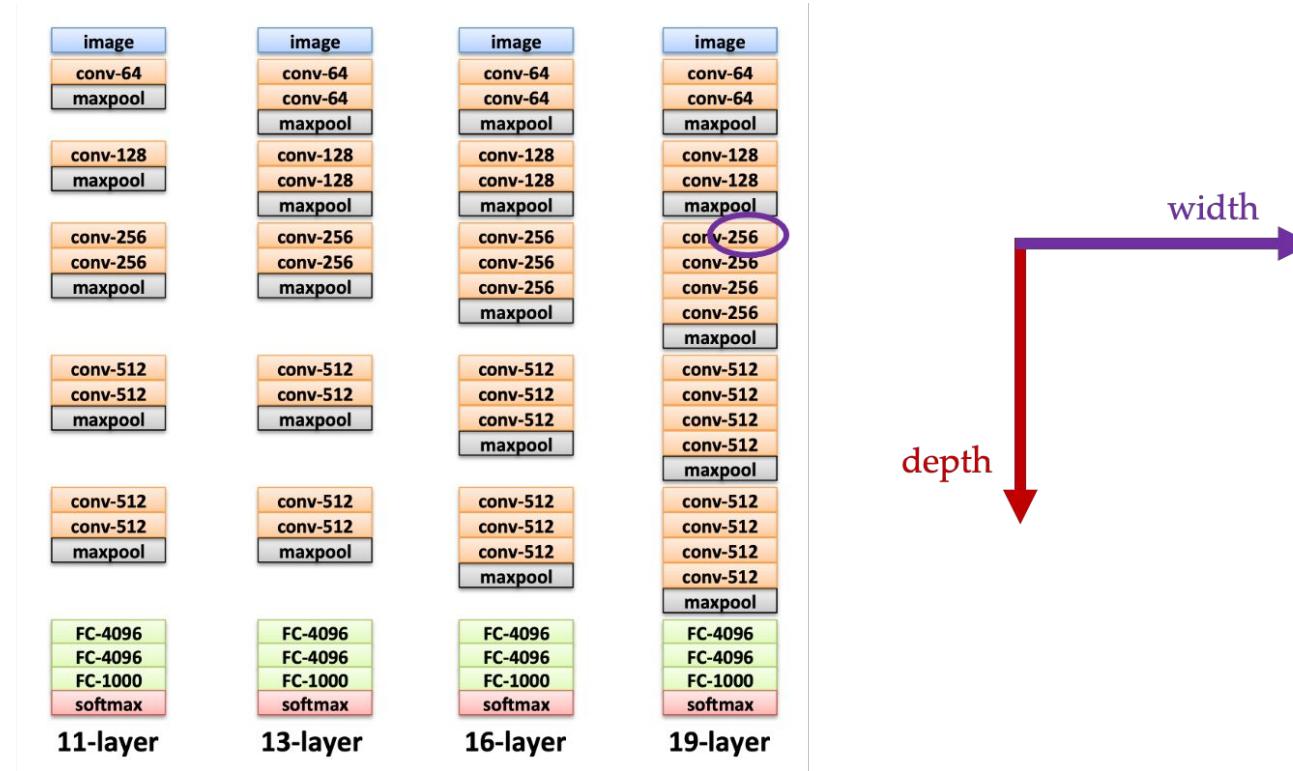
# VGGNet



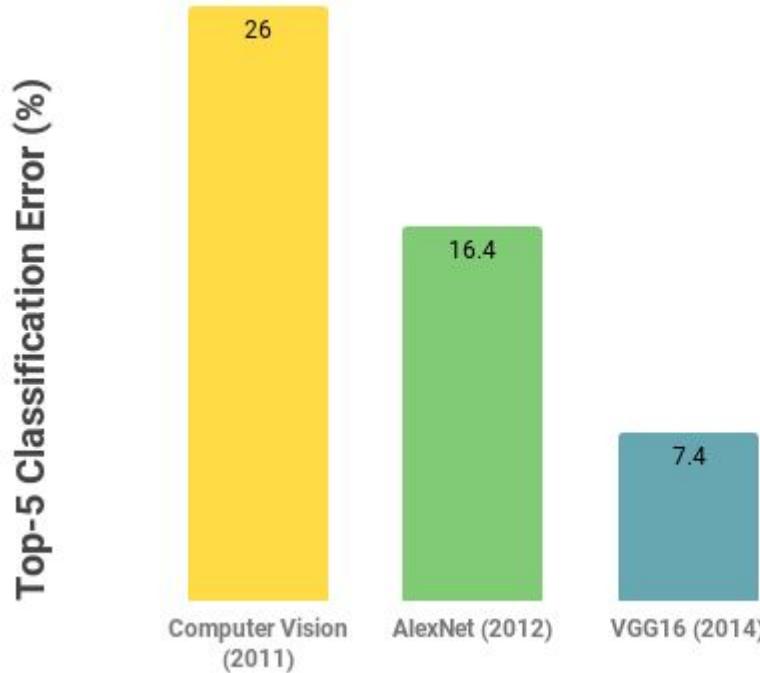
# VGGNet



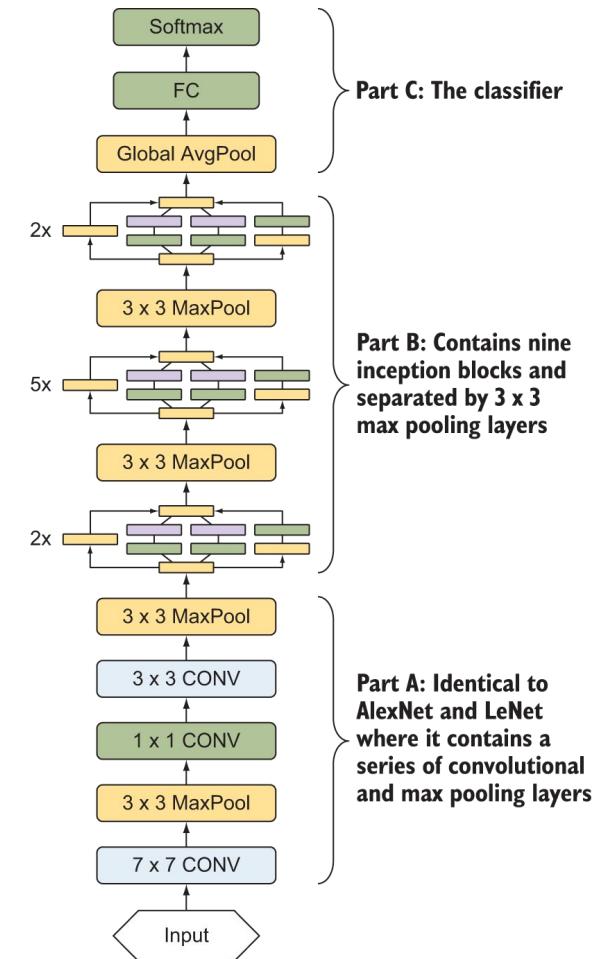
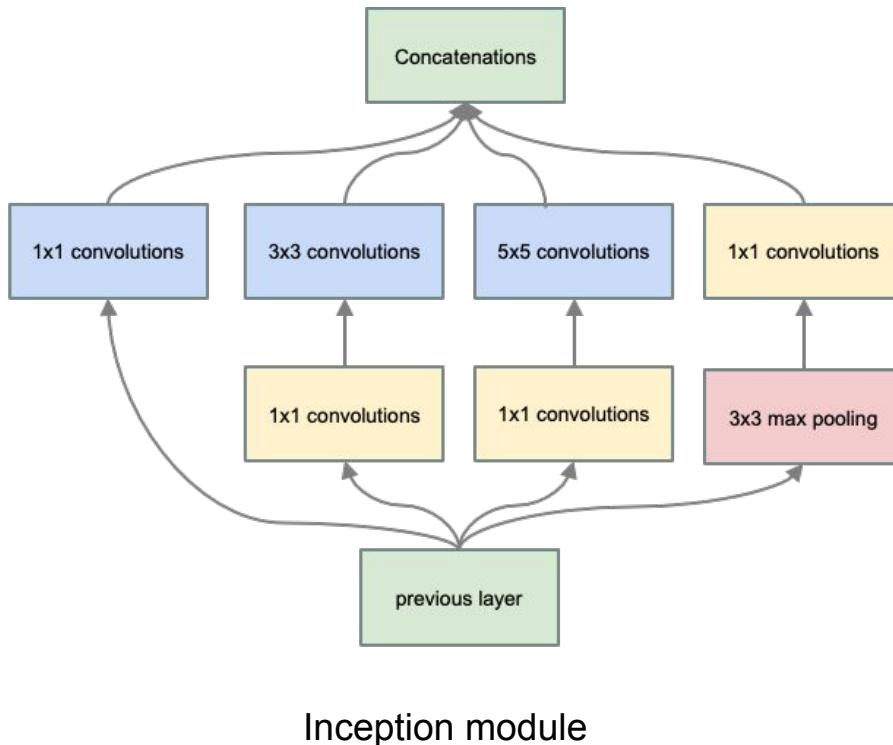
# VGGNet



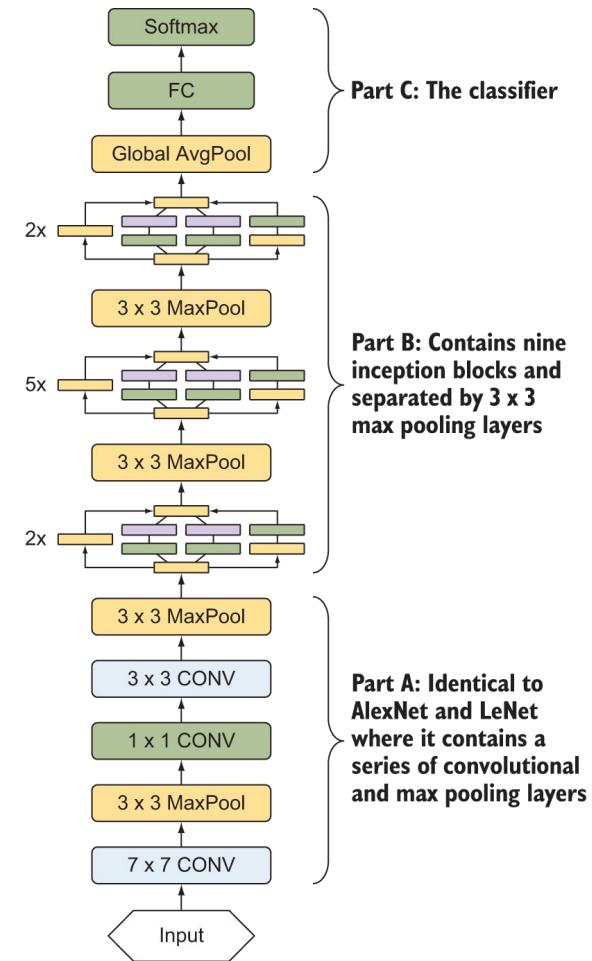
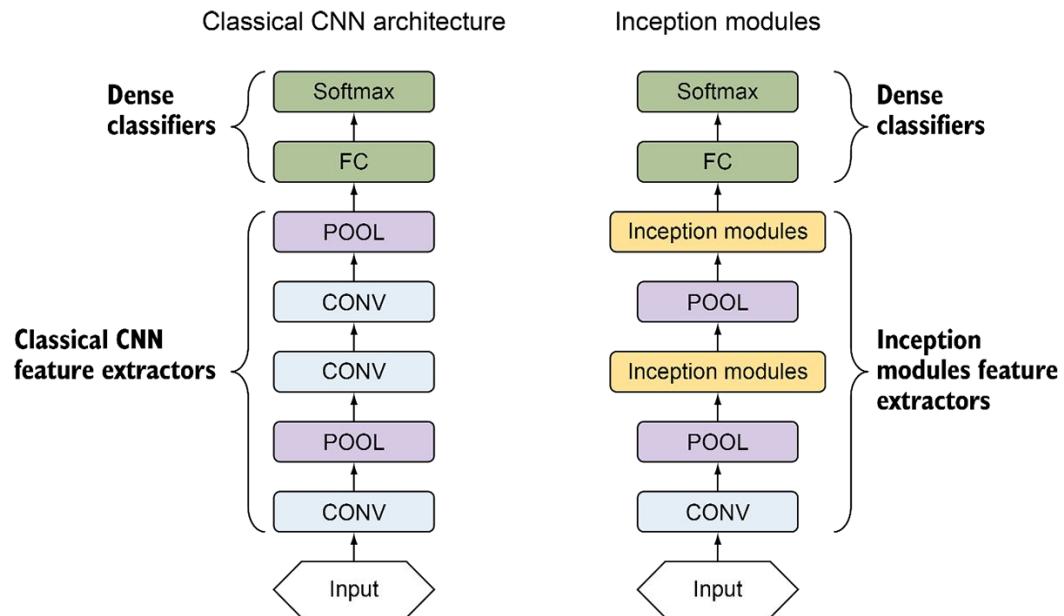
# The power of a curated dataset and competitions



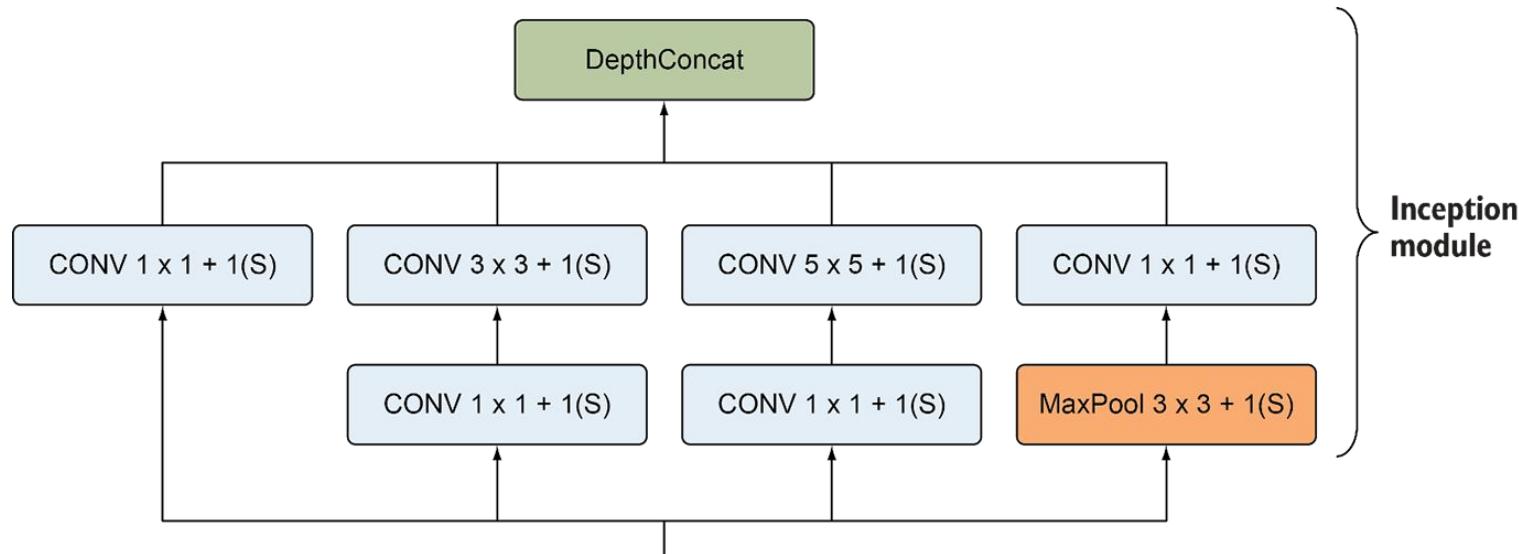
# GoogLeNet – Inception



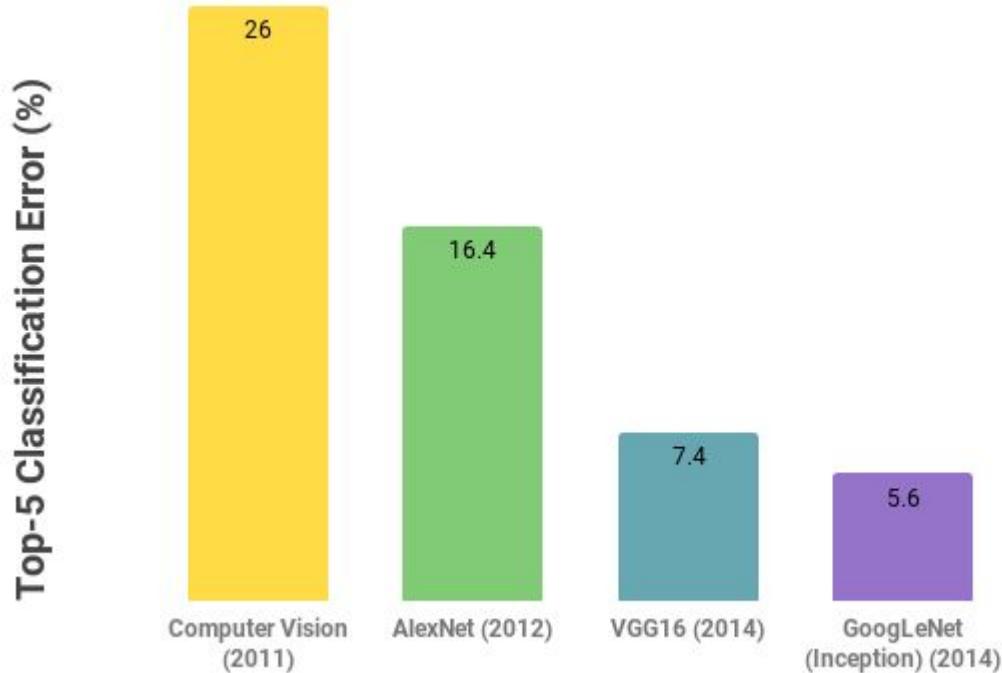
# GoogLeNet – Inception



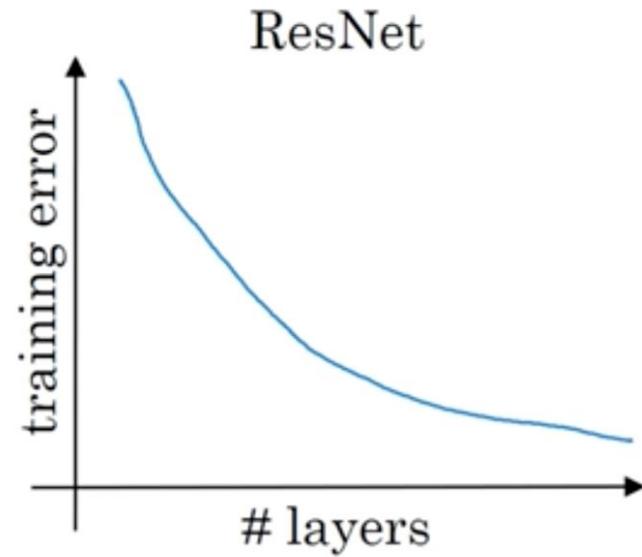
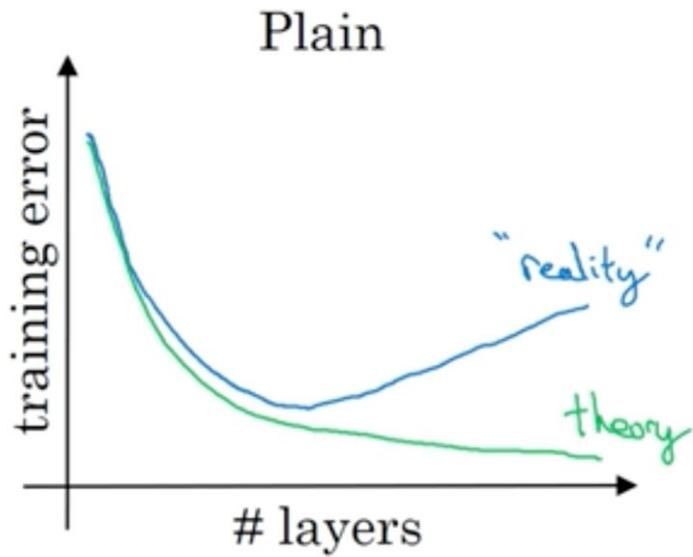
# GoogleLeNet – Inception



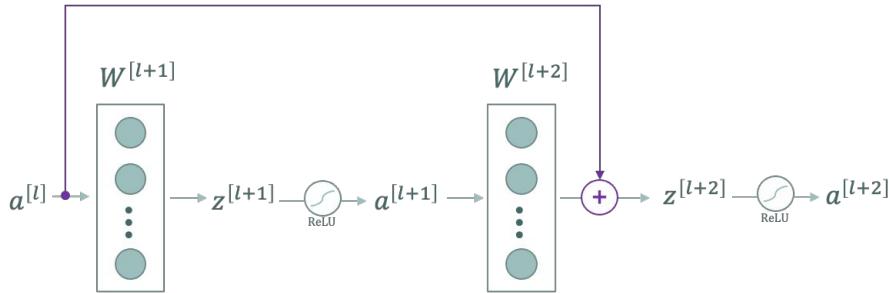
# The power of a curated dataset and competitions



# ResNet: the problem with deep layers



# ResNet: the problem with deep layers



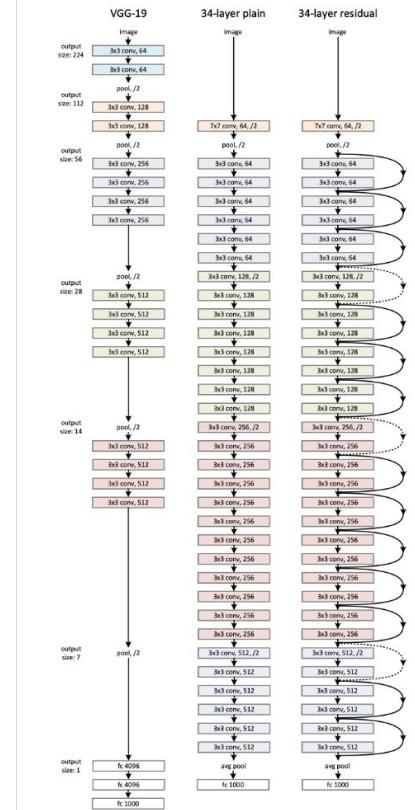
$$z^{[l+1]} = W^{[l+1]}a^{[l]} + b^{[l+1]}$$

$$a^{[l+1]} = \text{ReLU}(z^{[l+1]})$$

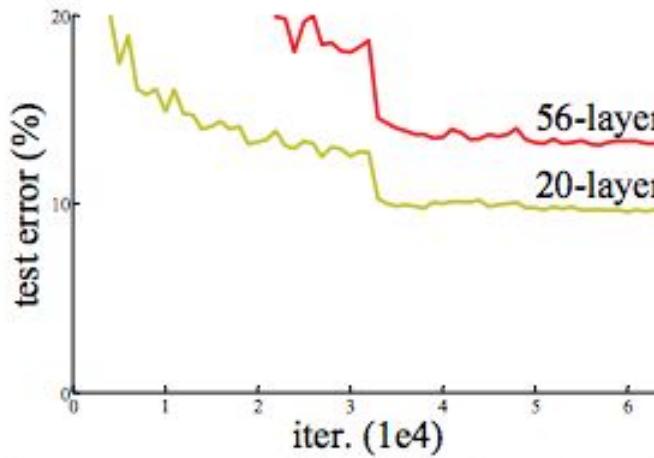
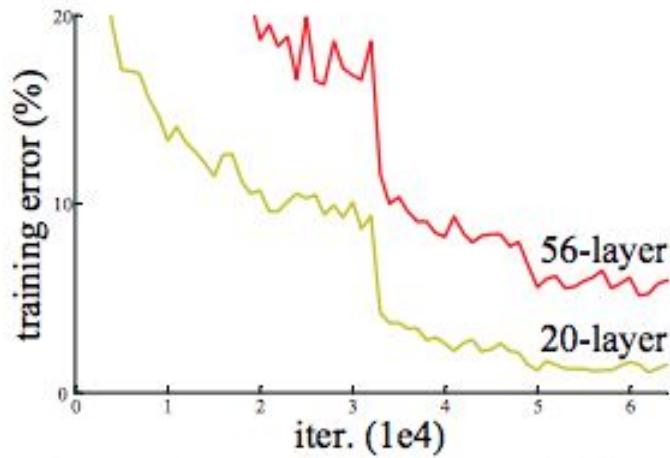
$$z^{[l+2]} = W^{[l+2]}a^{[l+1]} + b^{[l+2]}$$

$$a^{[l+2]} = \text{ReLU}(z^{[l+2]})$$

$$a^{[l+2]} = \text{ReLU}(z^{[l+2]} + a^{[l]})$$

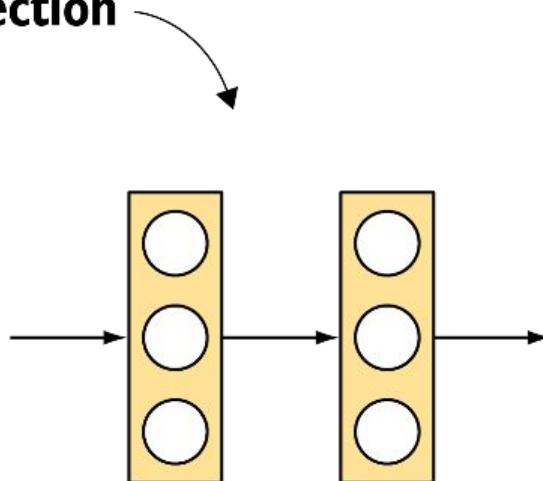


# ResNet: the problem with deep layers

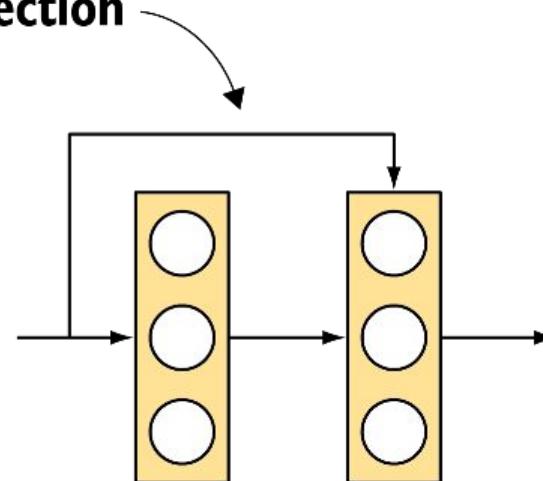


# ResNet: Residual networks and the idea of skip

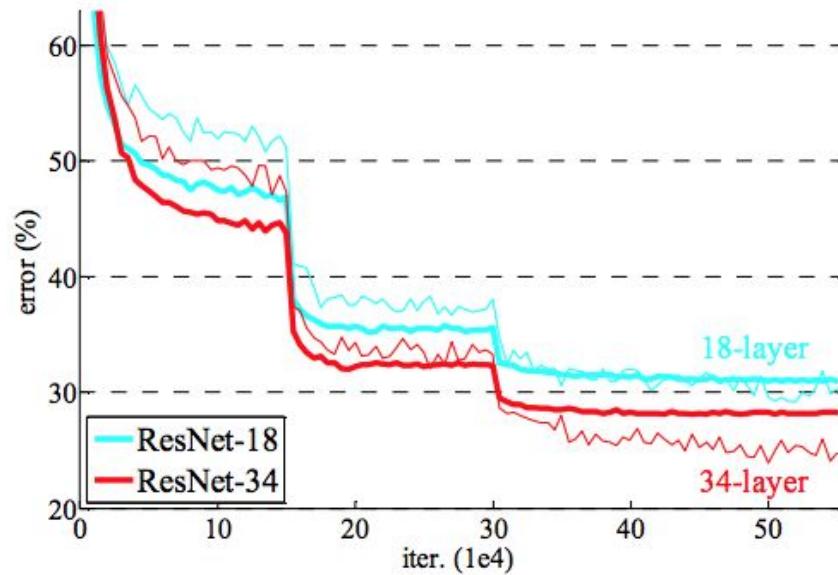
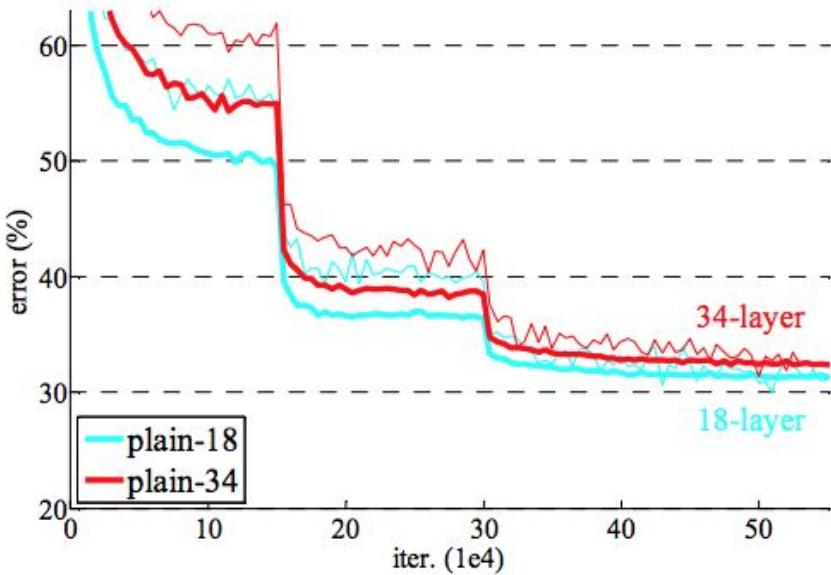
**Without skip connection**



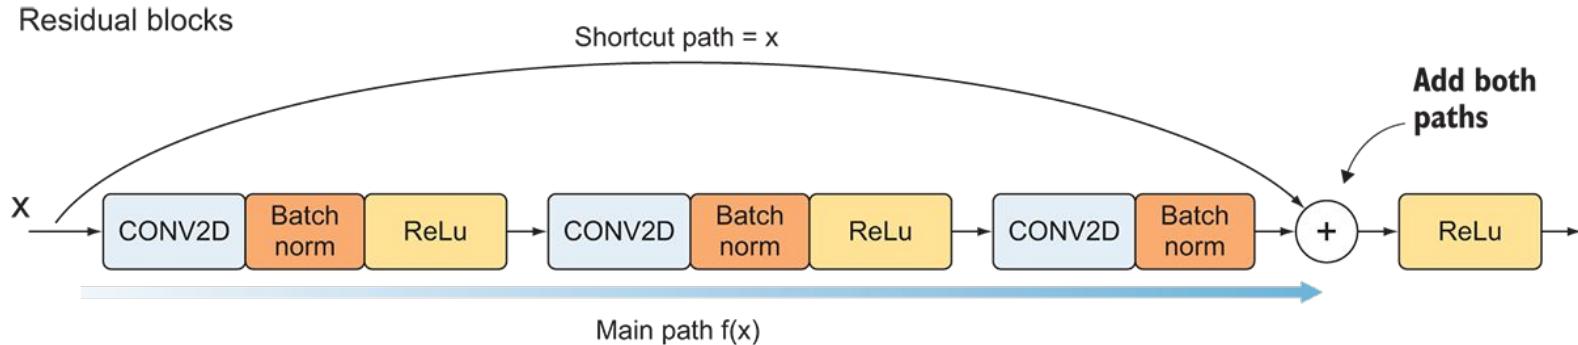
**With skip connection**



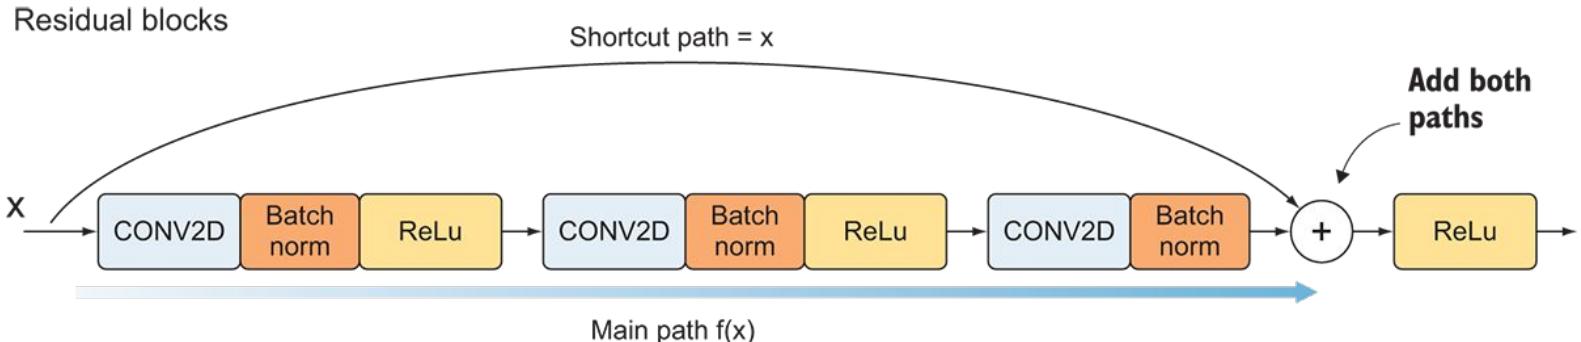
# ResNet: Residual networks and the idea of skip



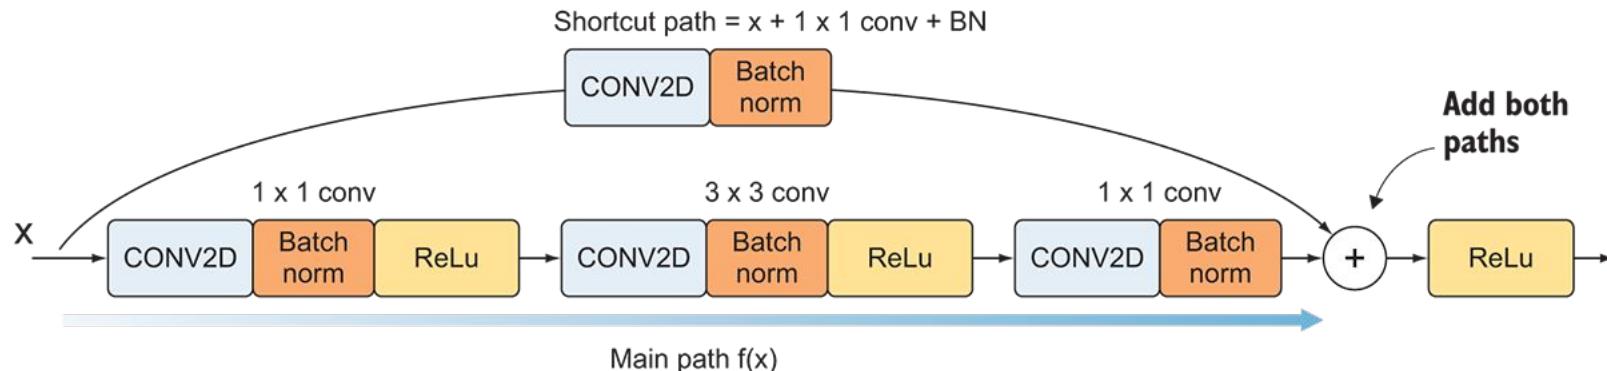
# ResNet: Residual networks and the idea of skip



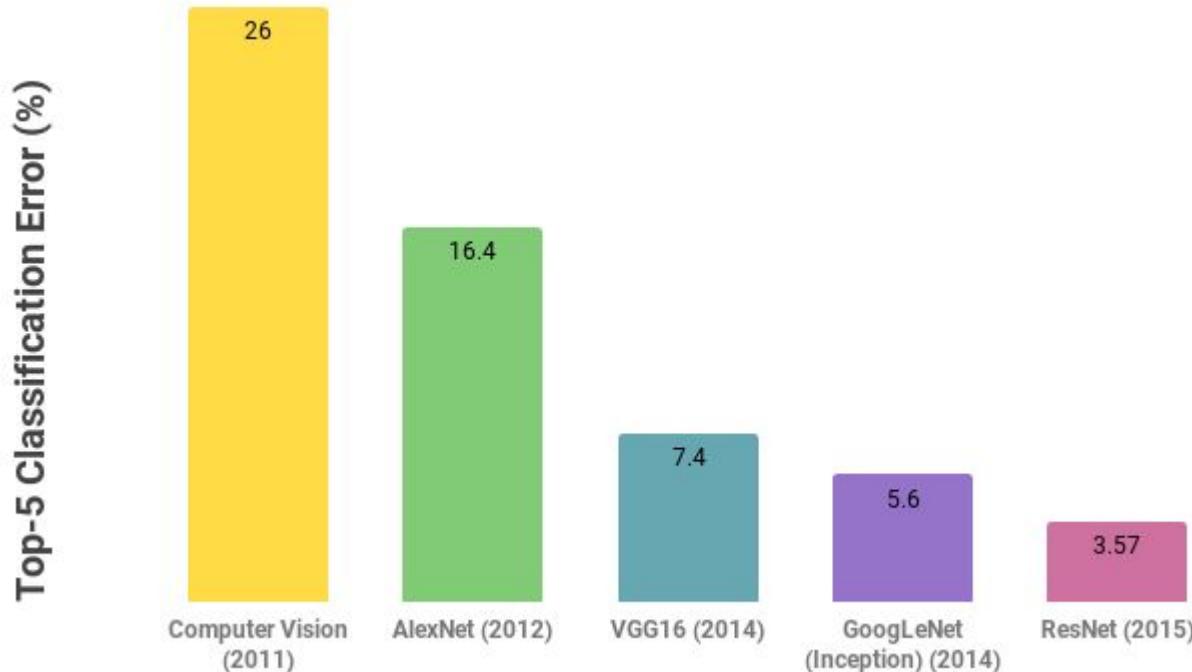
# ResNet: Residual networks and the idea of skip



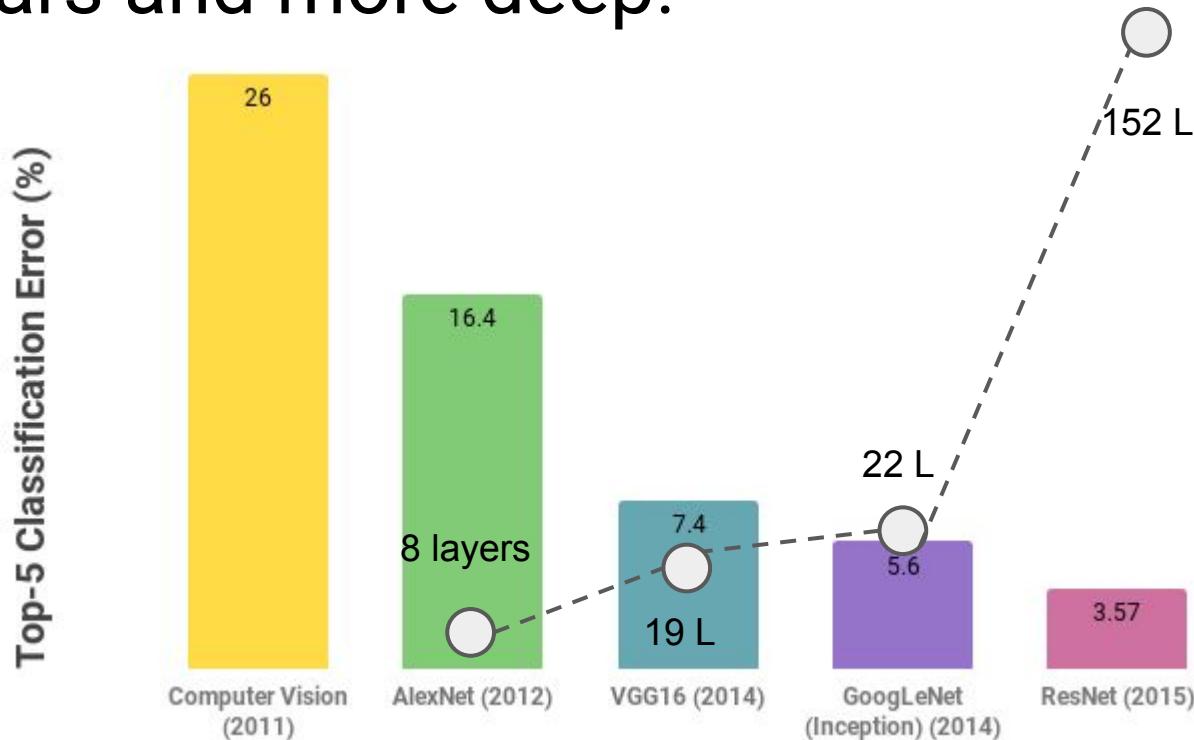
Bottleneck residual block with reduce shortcut



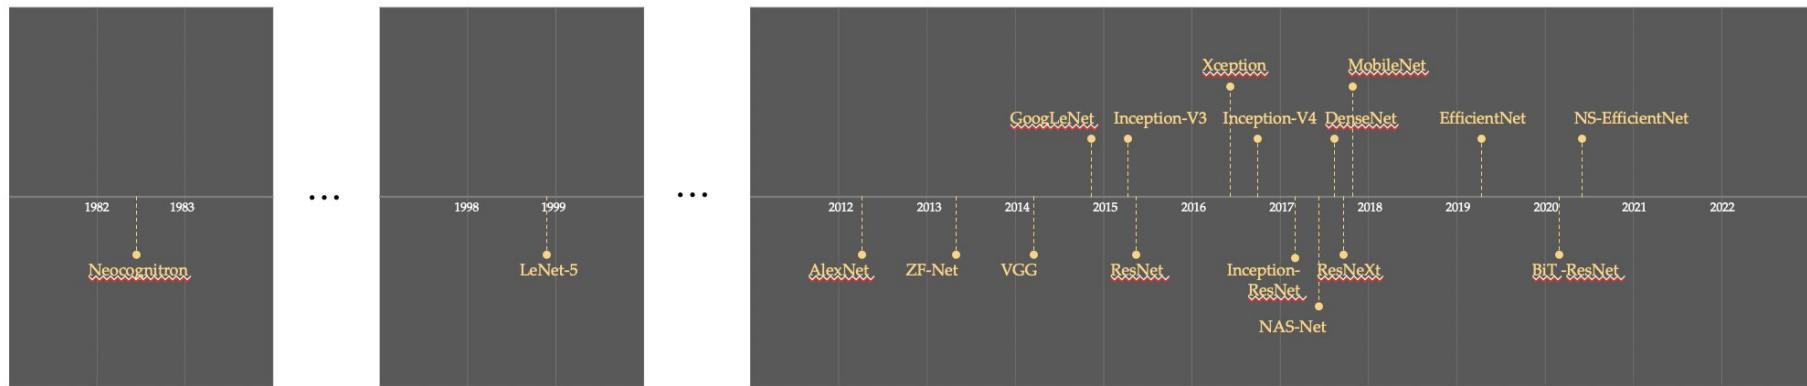
# The networks have become more powerful over the years



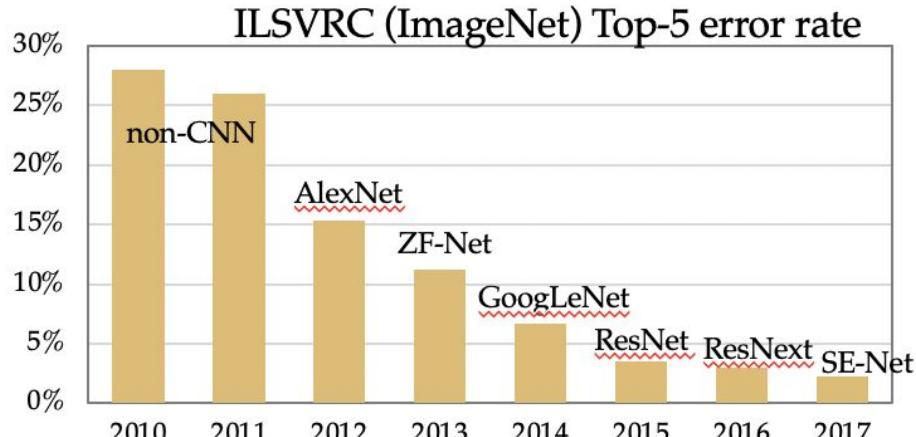
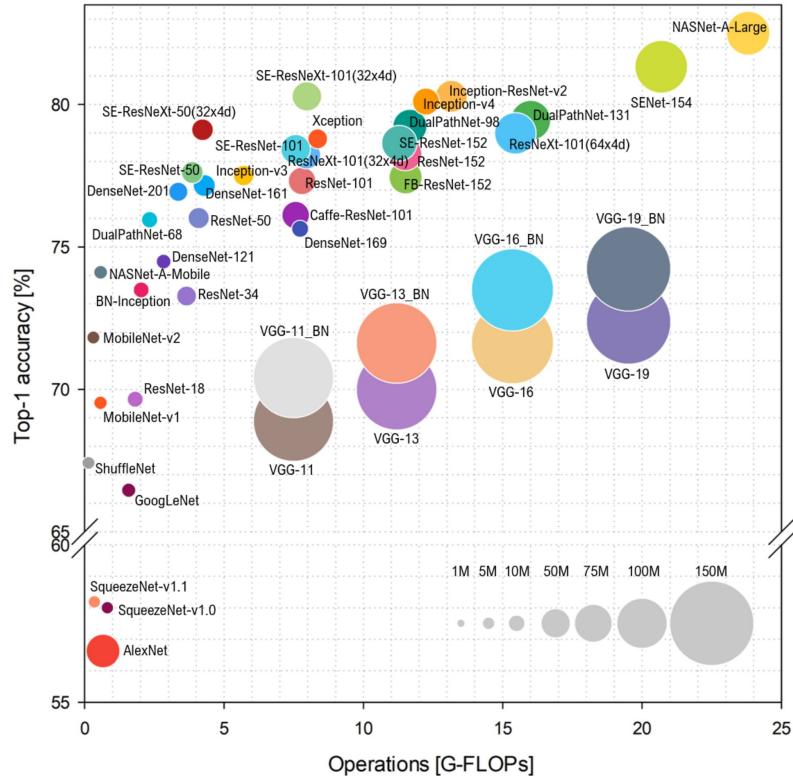
# The networks have become more powerful over the years and more deep!



# Timeline of CNN architectures

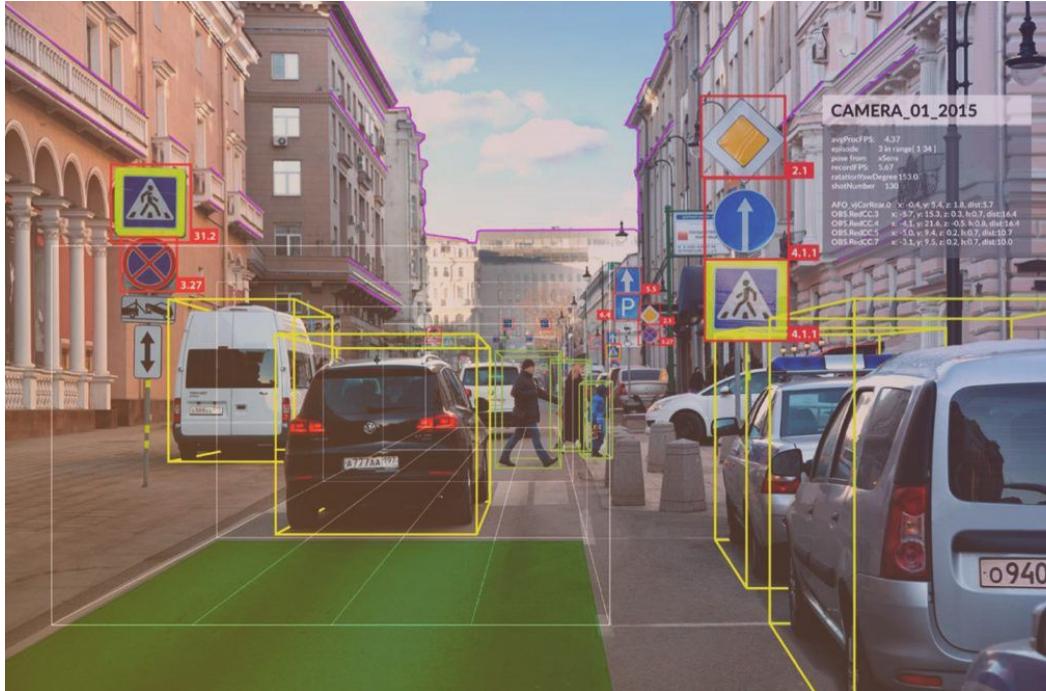


# Timeline of CNN architectures



# P5.

# CNN applications today!



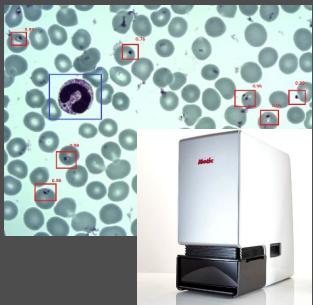
# Machine Learning for Pneumonia Detection on Lung Ultrasound

18 November 2021

Courosch Mehanian, Research Associate Professor

# Examples of Global Health projects

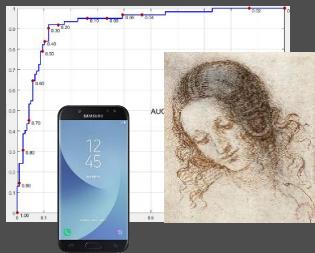
AI-powered  
Malaria  
microscope



AI-powered  
ultrasound  
for  
pneumonia  
detection



Automated  
Visual  
Evaluation for  
cervical  
precancer  
screening



Arktek for  
cold-chain  
transport



## Our mission

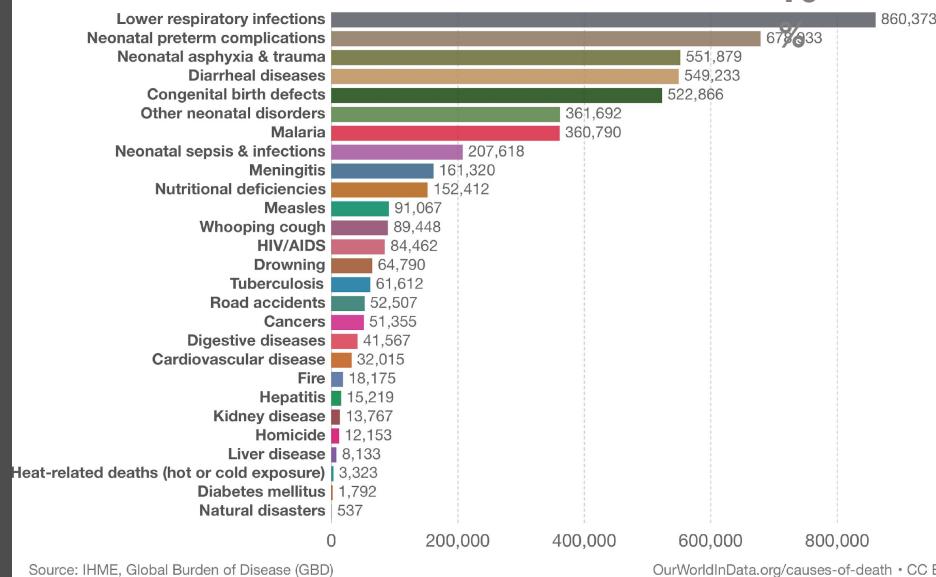
We build tools and technologies to address unmet needs in primary healthcare centers and last-mile service delivery in low- and middle-income countries (LMICs) around the world.

# Childhood pneumonia burden of disease

How can AI assist in early diagnosis of pneumonia in LMICs?

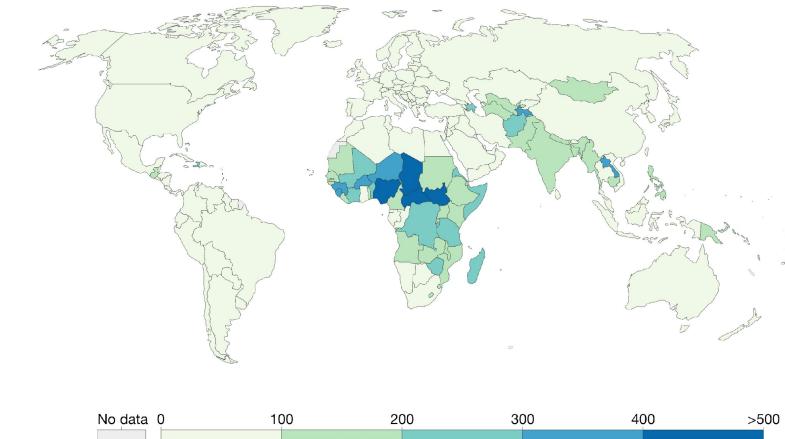
## Causes of death in children under 5, World, 2016

Annual number of deaths by leading causes in children under 5 years old.



## Death rates from pneumonia in children under 5, 2016

The annual number of deaths per 100,000 children under 5.



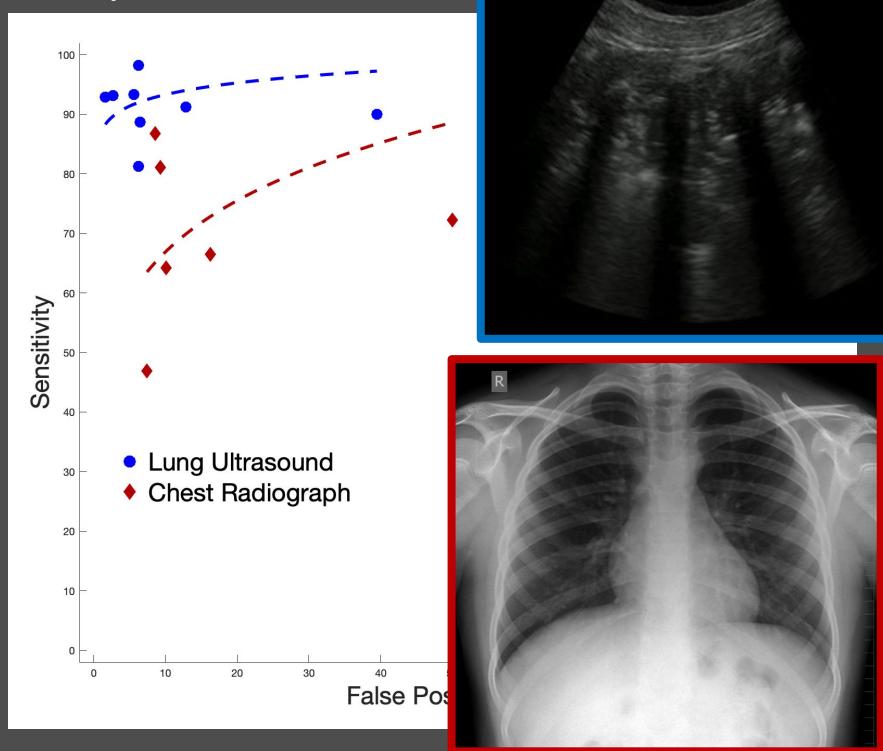
Source: IHME, Global Burden of Disease Study (2018)

Note: Deaths from 'clinical pneumonia', which refers to a diagnosis based on disease symptoms such as coughing and difficulty breathing and may include other lower respiratory diseases.

- 2,400 children per day died in 2016
- Most prevalent in South Asia and sub-Saharan Africa
- Can be detected and treated

# Lung ultrasound for pneumonia diagnosis?

Ultrasound features are artifact, not anatomy

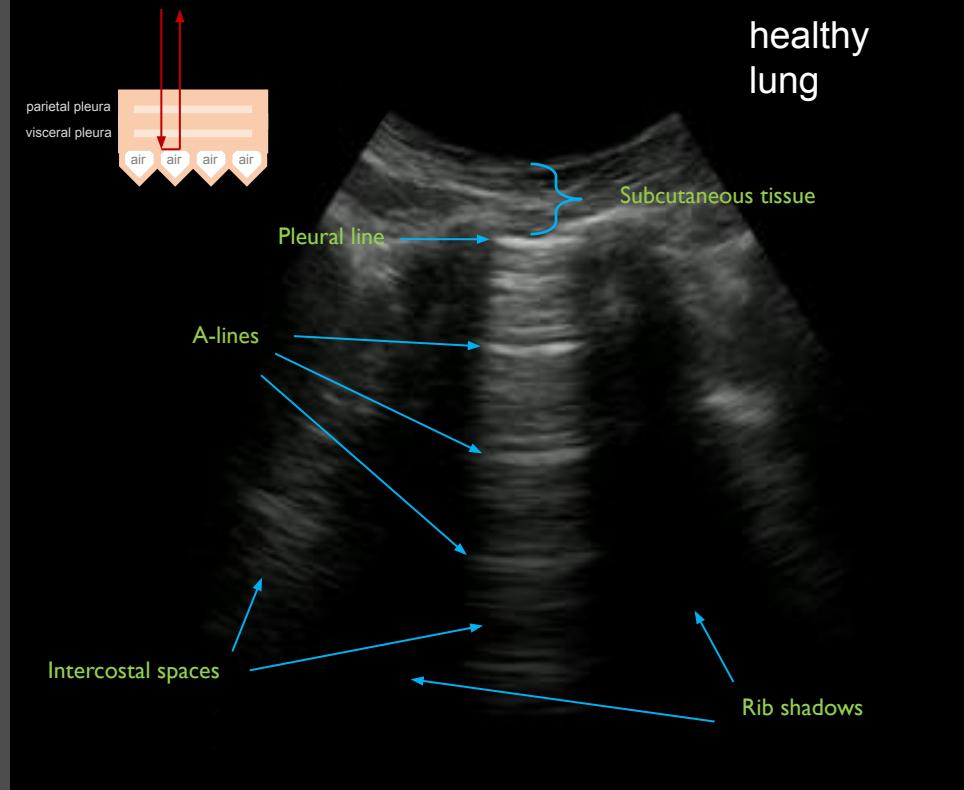
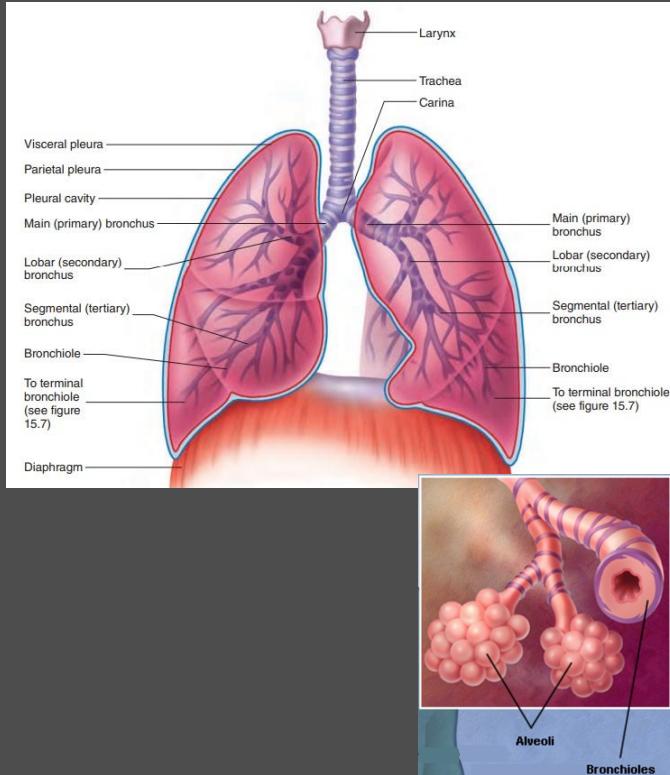


Advantages	LUS	CXR
Accuracy	good	medium
Safety	no radiation	ionizing radiation
Availability	point-of-care	needs referral
Portability	portable	usually fixed
Equipment cost	low	
Barriers	LUS	CXR
Acquisition skill	medium	low
Interpretation skill	difficult	medium
Ubiquity	not widely used	standard-of-care

*Lung ultrasound adoption could be enhanced by AI-assisted interpretation, especially in LMICs*

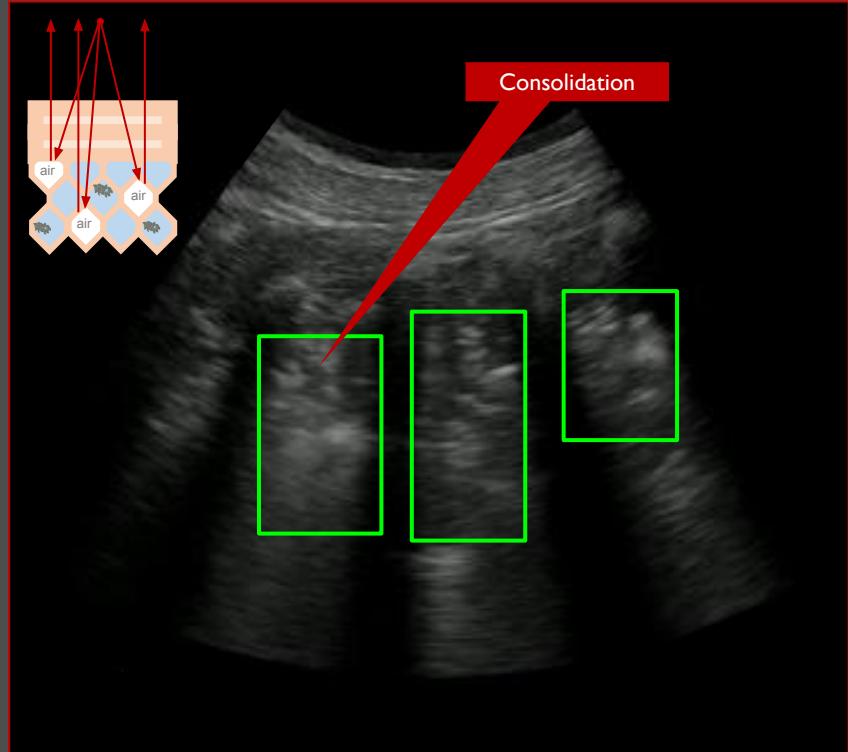
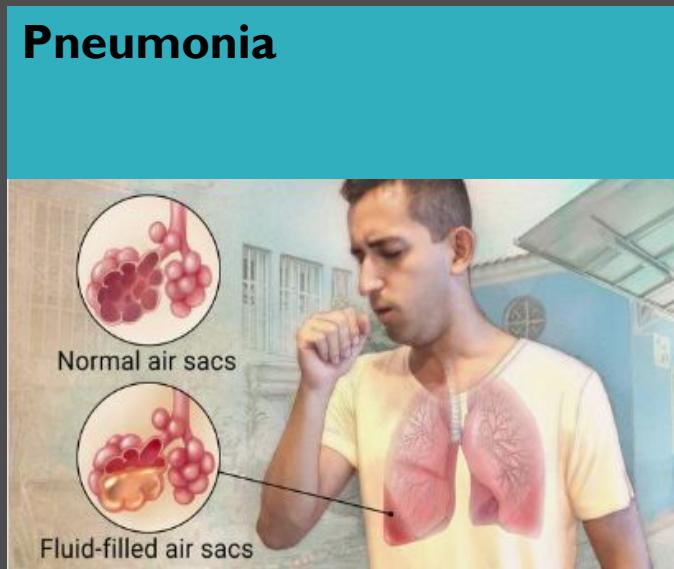
# Lung ultrasound 101

*Acoustic artifacts provide recognizable features*



# Typical pneumonia in lung ultrasound

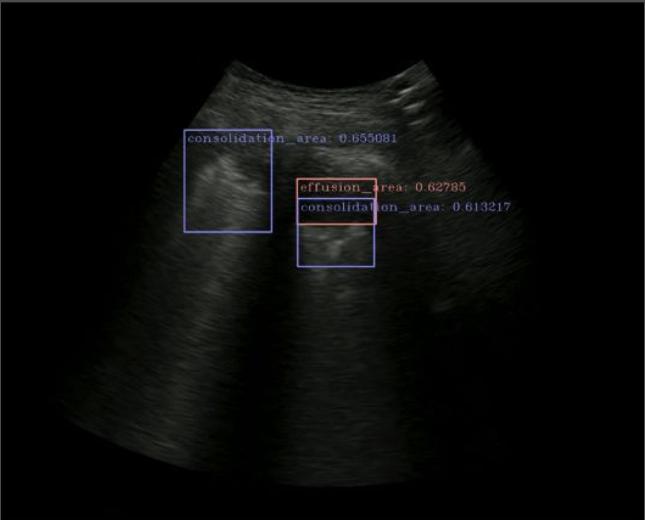
*Lung infection/inflammation partially fills lung air spaces with fluid*



# Use cases and requirements

- Detect consolidation and pleural effusion
- Adults ( $\geq 18$  years old)
- Children ( $\geq 10$  kg and  $< 18$  years old)
- Body Mass Index (BMI  $\leq 40$  kg/m<sup>2</sup>)
- Sensitivity  $\geq 90\%$  at video level
- Specificity  $\geq 90\%$  at video level
- Draw bounding boxes around detections
- Time-to-results  $\leq 1$  minute, post-exam on Samsung S4 tablet
- Data collected by sonographer trained in lung ultrasound

- Patient with pneumonia-like symptoms
- Evaluation includes AI-assisted ultrasound plus other clinical measurements
- Patient is released, prescribed antibiotics, or referred for follow-up



# Data collection

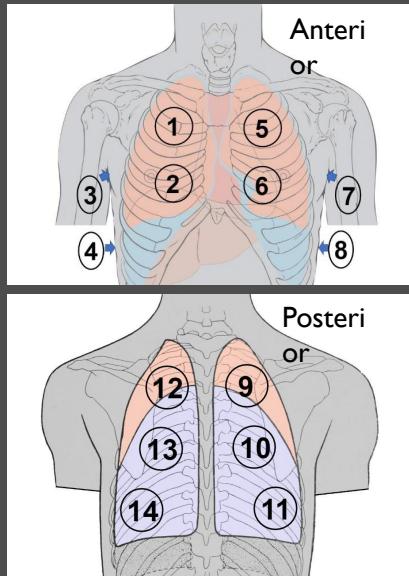
## Image Acquisition



## Image Annotation



## Algorithm Development



Oregon Health &  
Science University

OHSU Doernbecker  
Children's Hospital

Columbia University  
Medical Center

MedStar Hospital  
Washington, DC

Brooke Army  
Medical Center,  
San Antonio, TX

Womack Army  
Medical Center,  
Fort Bragg, NC

Tripler Army Medical  
Center, Honolulu, HI

	Children	Adults
Patients	306	617
Ultrasound exams	320	1,225
Annotated videos	1,726	6,498
Annotated frames	103,560	389,880

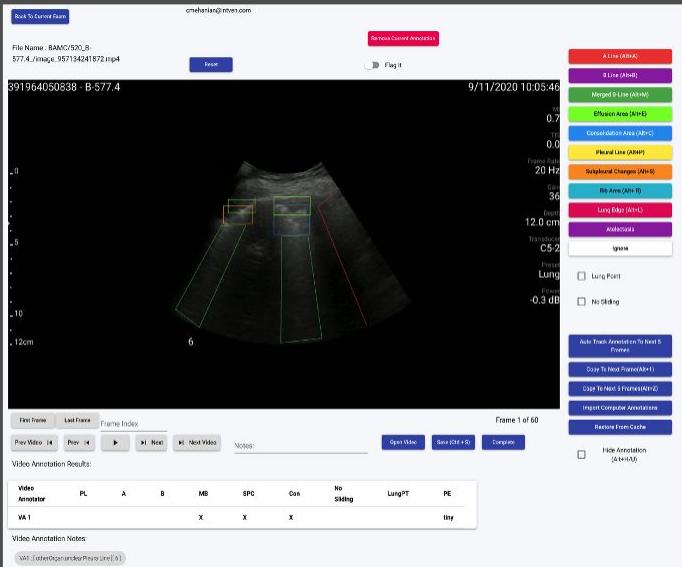
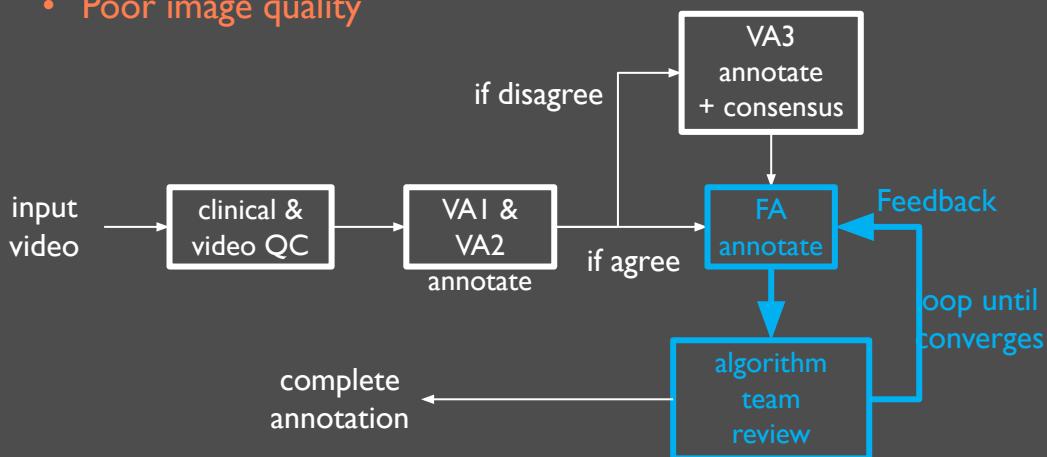
\*Volpicelli, et al., International evidence-based recommendations for point-of-care lung ultrasound, *Intensive Care Medicine* (2012), 38:577-591

# Curating high-quality annotated image library



## Video exclusion criteria:

- Incorrect acquisition parameters
- Excessive transducer movement
- Incorrect transducer positioning
- Poor image quality



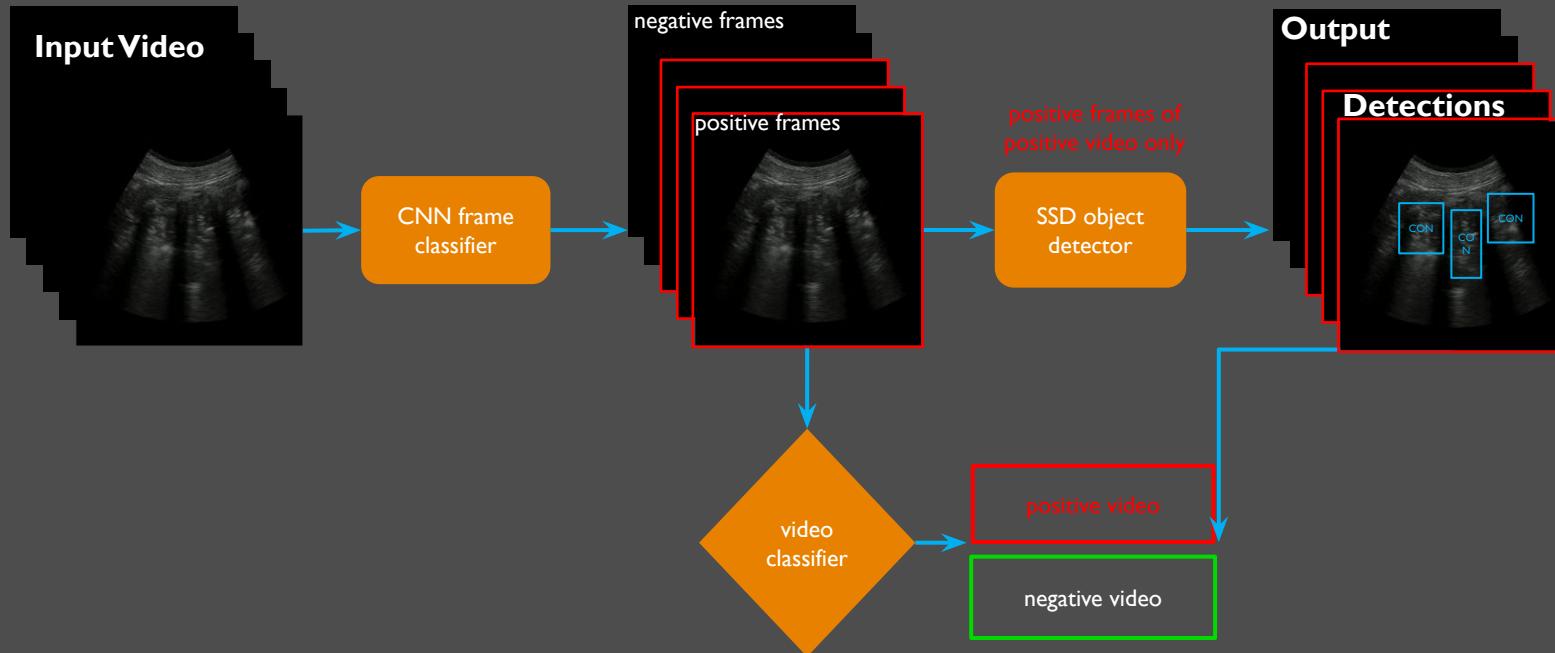
Careful frame-by-frame annotation by expert radiologist

# Cascade architecture for classification & localization

Image Acquisition

Image Annotation

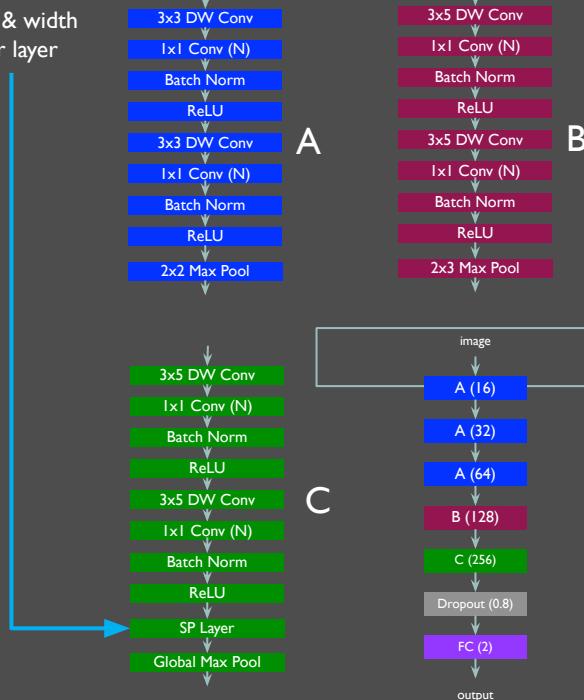
Algorithm Development



# Classification and detection models

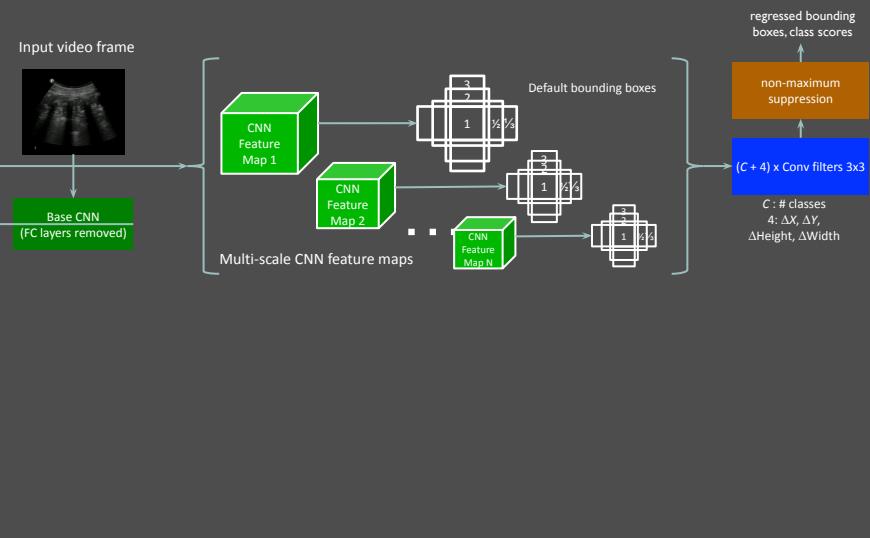
## Adult consolidation classification

- VGG-like CNN architecture
- Simplify convolutions
- Optimize depth & width
- Add spatial prior layer



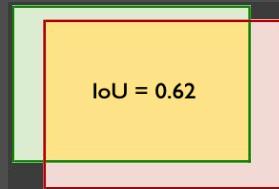
## Detection models (all lung features)

- SSD architecture is accurate
- Features at multiple scales, aspect ratios
- Efficient, suitable for on-device inference



# Summary of results

Specifications are close to being met on holdout sets



## ADULT CONSOLIDATION

	patients	exams	neg. videos	pos. videos
TRAIN	46	102	984	358
VAL	15	21	55	66
HOLDOUT	26	31	132	132
	sensitivity		specificity	avg. IoU
HOLDOUT	0.88		0.90	0.60

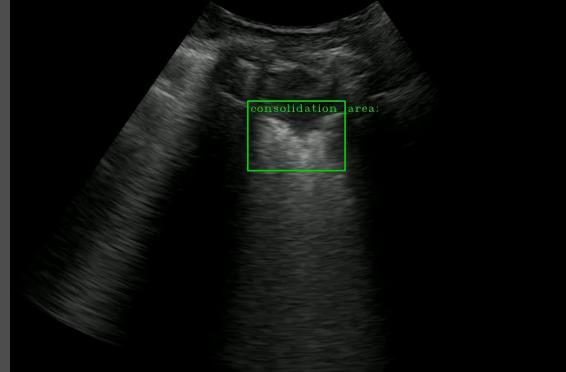
Healthy lung



## PEDIATRIC CONSOLIDATION

	patients	exams	neg. videos	pos. videos
TRAIN	55	61	466	402
VAL	26	31	178	136
HOLDOUT	27	27	127	133
	sensitivity		specificity	avg. IoU
HOLDOUT	0.88		0.89	0.62

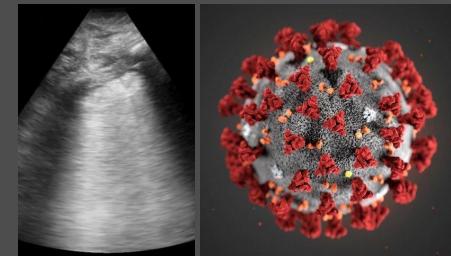
Consolidation



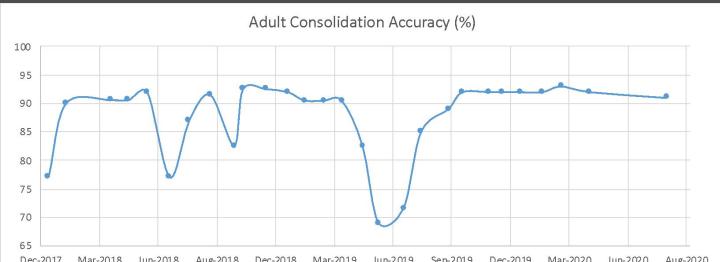
All models execute  
in < 30 sec/video

Articles
AAAS 2018
ASA 2018
MHSRS 2018 * 2
MICCAI 2018
MHSRS 2019 * 2
MICCAI 2019
MHSRS 2020 * 2
Critical Care 2019

# Summary and Conclusions



- Developed algorithm for AI-assisted detection of pneumonia with lung ultrasound
  - Worked with a global leader in ultrasound technology
  - Algorithms may be directly applicable to smoke inhalation, COVID-19 and other ultrasound applications are in the pipeline: compartment syndrome, FAST, optic nerve-sheath diameter
- Learned
  - Importance of large, high-quality, diverse data set
  - Importance of high-quality annotations suitable for machine learning
  - Importance of collaboration between doctors and ML scientists
- Learn-as-you-go



# Acknowledgements

*Funded by Bill and Melinda Gates Foundation Trust and DARPA*

Sourabh Kulhare	Michele Griffin	Kenton Gregory	Todd Graham	Matthew Hepburn
Rachel Millin	Bruce Goldstein	Cynthia Gregory	Katelyn Hostetler	Matthew Kappes
Dipayan Banik	Davina Inslee	Meihua Zhu	Annie Cao	Paul Sheehan
Liming Hu	Doug Bradley	Hua Xie	Michael Stone	Andrea Cantu
Charles Delahunt	Sheana Creighton	Fen Wang	Andrew Hersh	Jeffrey Shupp
Zohreh Laverriere	Debra Peat	Xijun Zhang	David Kessler	Laura Johnson
Matthew Horning	Megan Bettilyon	James Jones	Cristian Madar	Peter Weimersheimer
Benjamin Wilson	Shannon Kuyper	Jack Lazar	Damon Forbes	Kristin Dwyer
Xinliang Zheng	Renée Rivers	Amber Halse	Brian Shahan	Jacob Avila

# Summary of lessons learned

Worked well ...	Lung ultrasound
Curation of high-quality, diverse, multisource data	✓
Feedback mechanism for establishing ground truth	✓
Collaboration between ML people and SMEs	OHSU Clinicians
Focus on product, patient, and use-case	Video-level localization and classification
Partnership with commercial entity	Ultrasound Co.
Early engagement with regulatory agency	U.S. FDA
Keep an eye on ...	
Explore vs. exploit: Bayesian approach	Strike the right balance, hedge your bets
Bottom-up vs. top-down: Collective responsibility	Joint ownership of goals and timelines

# Questions ?

# Next lecture: *CNN in action*

